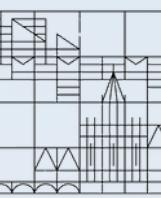


# 02 | Introduction to Neural Networks

Giordano De Marzo

<https://giordano-demarzo.github.io/>

Deep Learning for Social Sciences



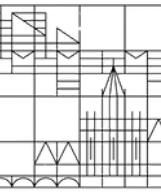
# Outline

1. Basic Concepts and Notation

2. The Perceptron

3. Limits of the Perceptron

4. Shallow Neural Networks



# Basic Concepts and Notation



# Learning from Data

Machine learning represents a paradigm shift in how computers solve problems.

**Traditional programming:** Humans write explicit rules for the computer to follow

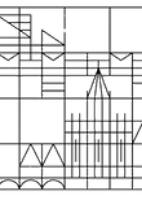
**Machine learning:** Computers discover patterns and rules from examples in data

This data-driven approach can be categorized into **two main types**:

- **Unsupervised learning:** Finding patterns in unlabeled data
- **Supervised learning:** Learning from labeled examples to make predictions

Supervised learning further divides into:

- **Classification:** Predicting categories (spam/not spam, dog/cat)
- **Regression:** Predicting numerical values (price, temperature)



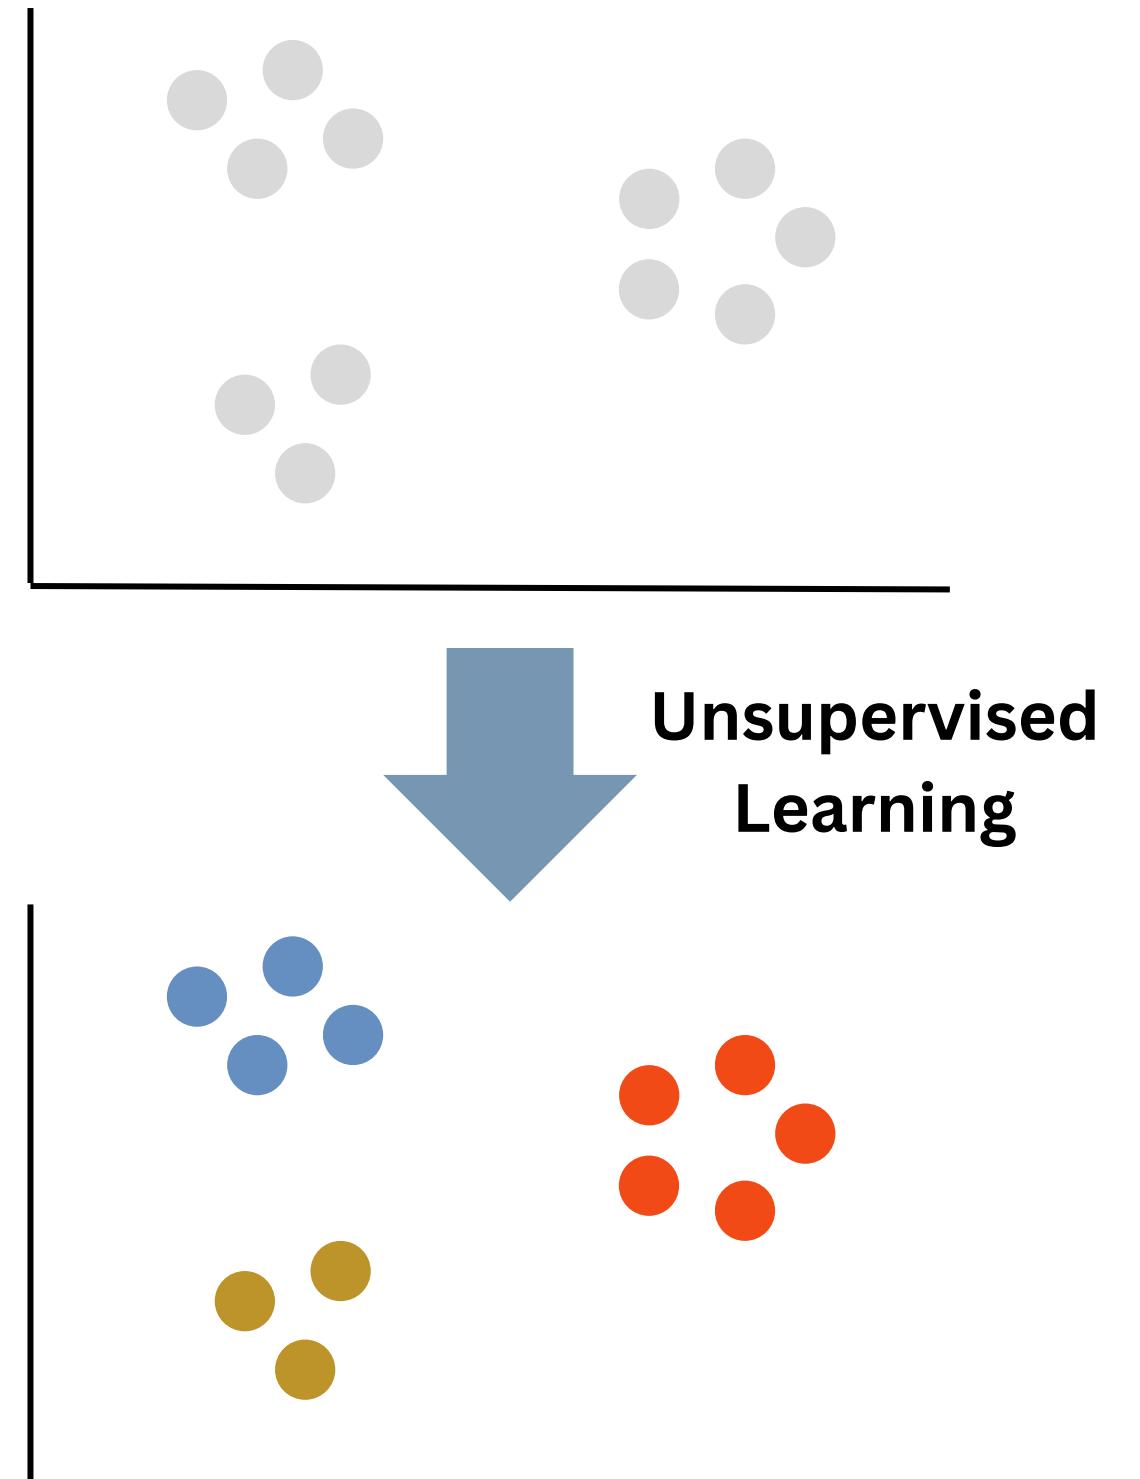
# Unsupervised Learning

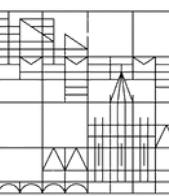
Unsupervised learning algorithm find structure in data without explicit guidance

- No "correct answers" or labels are provided in training
- The algorithm must discover meaningful patterns
- Unlabeled data is typically more abundant

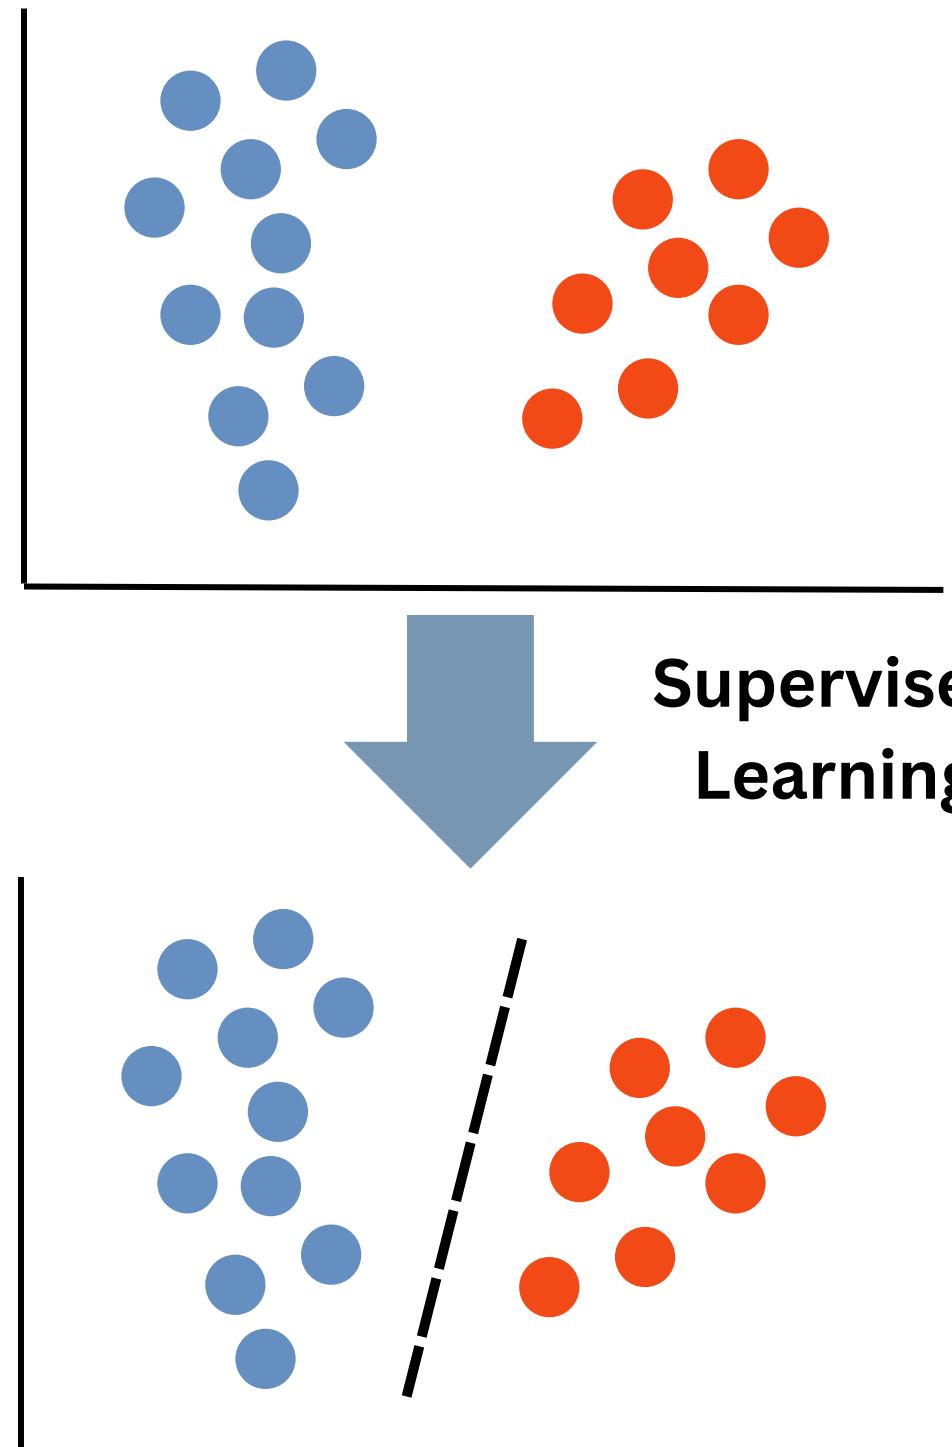
Several distinct tasks fall under the unsupervised learning umbrella.

- Clustering algorithms
- Dimensionality reduction techniques
- Anomaly detection methods





# Supervised Learning



Supervised learning relies on examples where the correct answer is already known.

- The fundamental idea is to learn from input-output pairs
- These paired examples serve as a teacher, showing the model what output should result from each input
- Creating these teaching examples typically requires human effort

The essential goal is to generalize beyond the training examples.



# Regression and Classification

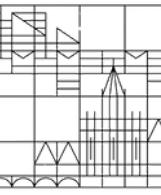
Supervised learning problems generally fall into two categories

## Regression tasks

- The output is a continuous numerical value representing a quantity or magnitude.
- Common applications include predicting house prices based on features, forecasting temperatures, or estimating a person's age from a photograph.

## Classification tasks

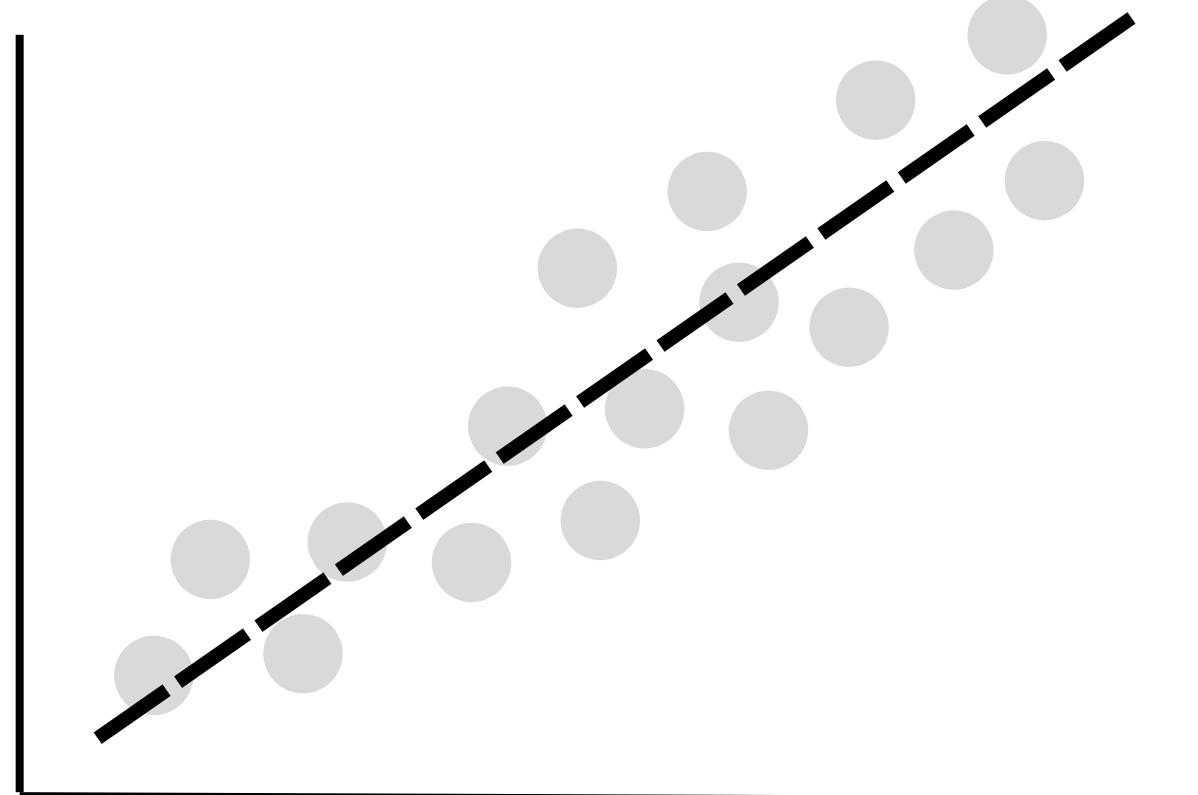
- The goal is to assign inputs to distinct categories. The output is either a discrete class label or a probability distribution across possible classes.
- Examples include filtering spam emails, diagnosing diseases from symptoms, or recognizing handwritten digits.



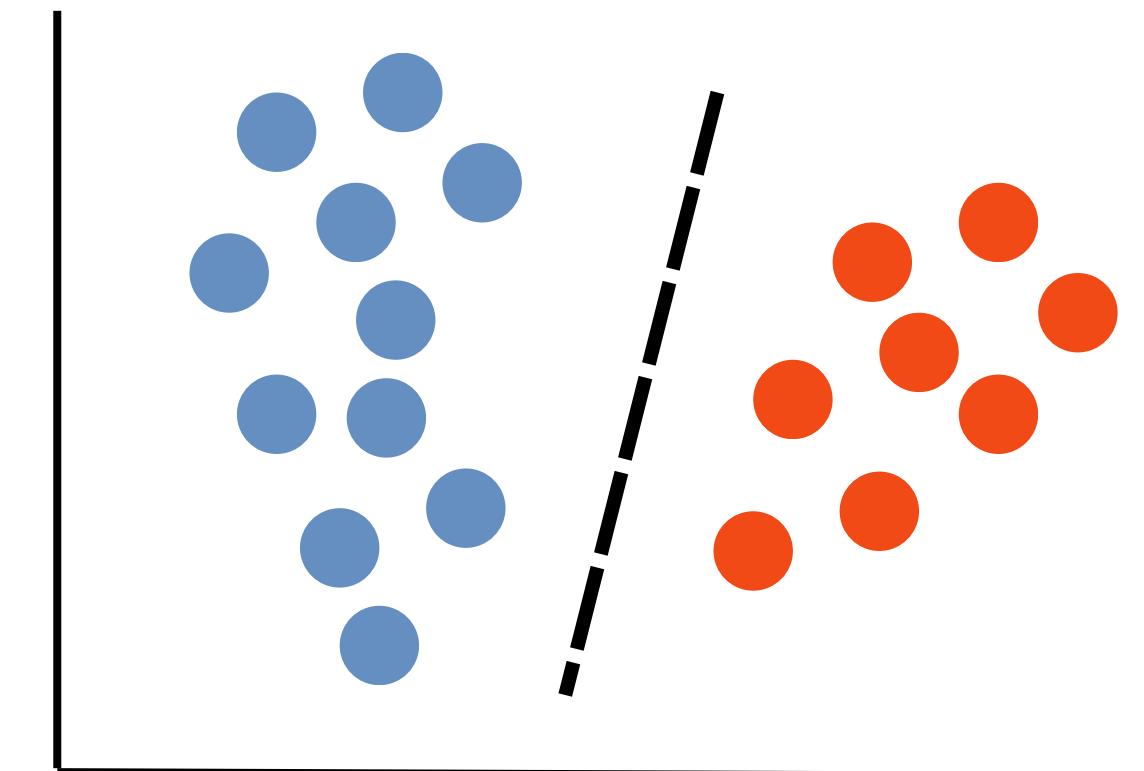
# Regression and Classification

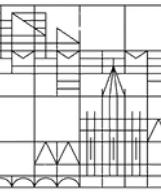
Supervised learning problems generally fall into two categories

**Regression**



**Classification**

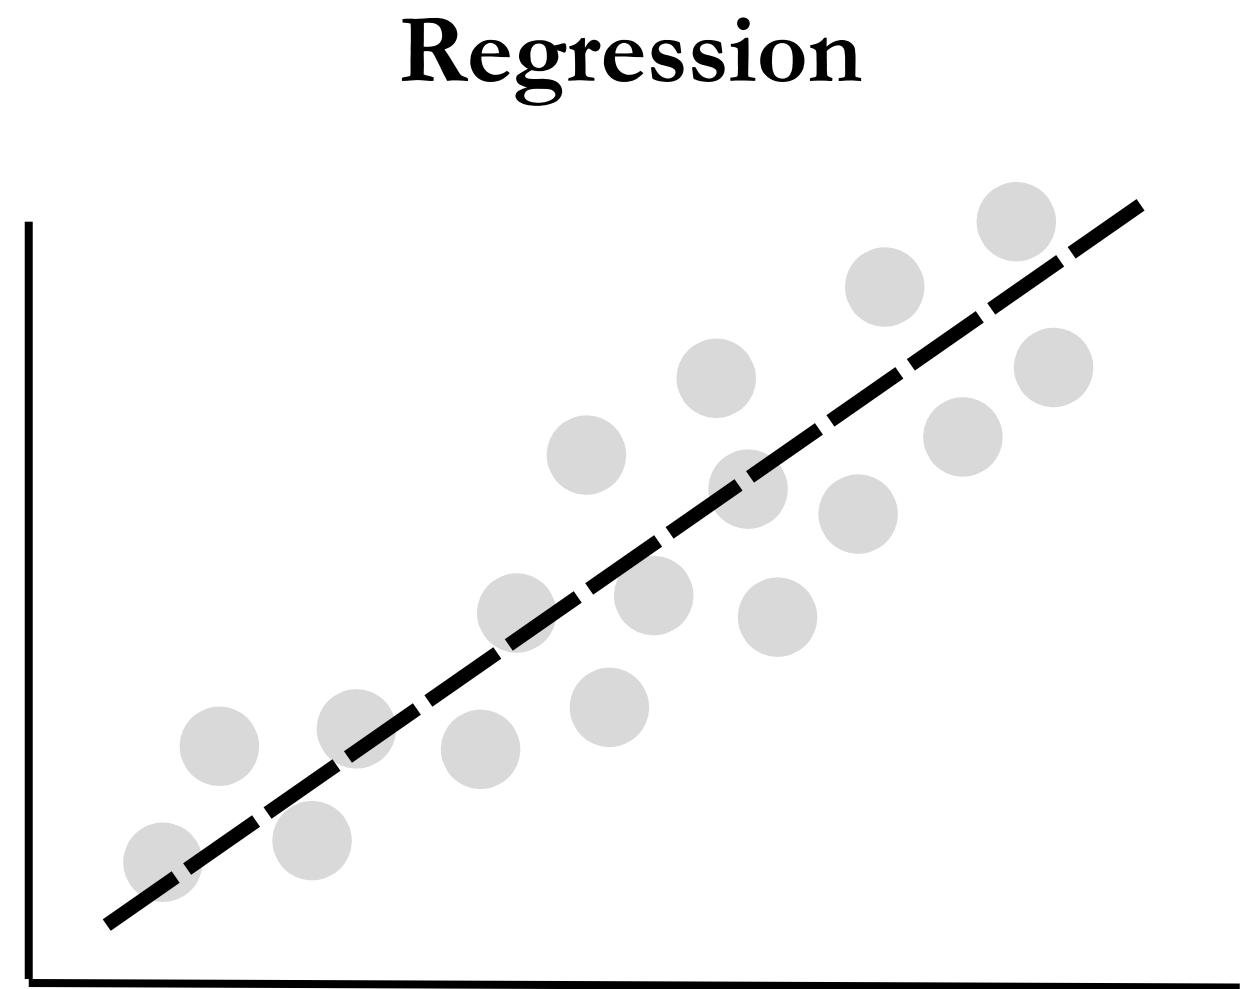


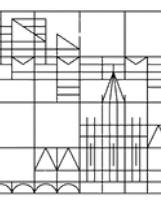


# Example: Linear Regression

Linear regression is the most simple example of supervised learning

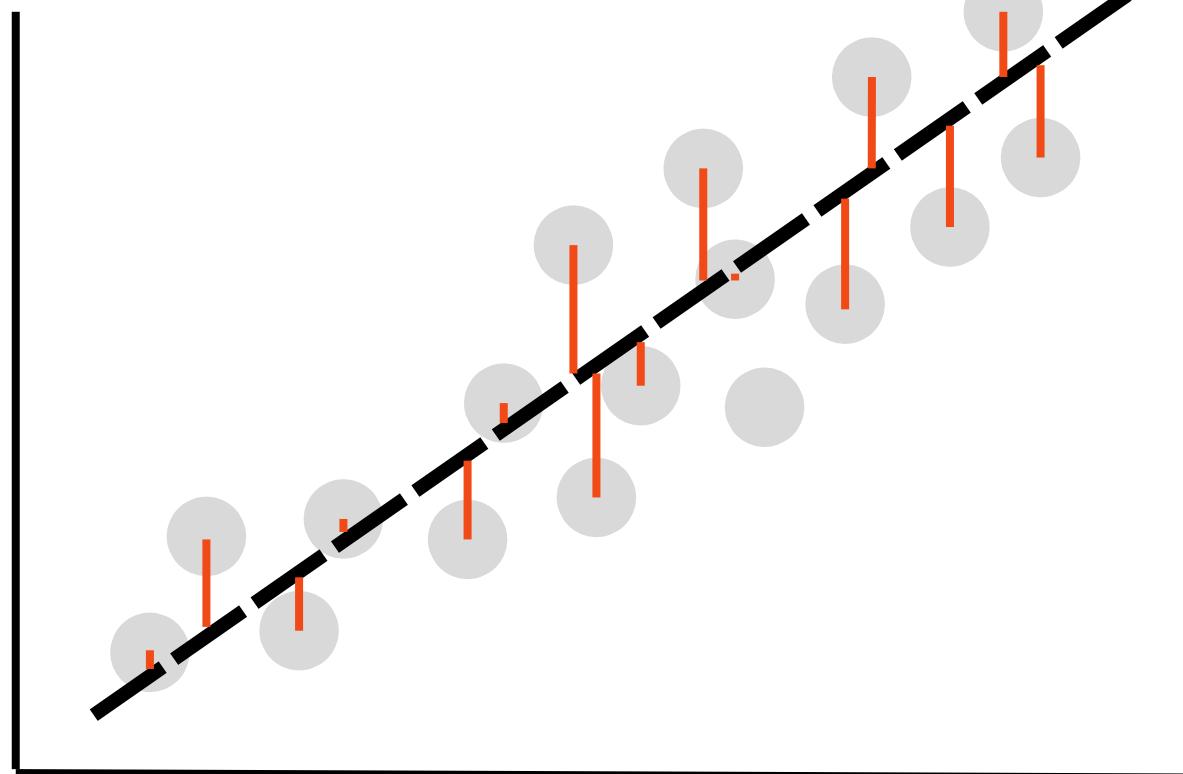
- The model is simply:  $y = wx + b$ 
  - $y$  is the output (prediction)
  - $x$  is the input feature
  - $w$  is the weight, and  $b$  is the bias.
- We provide training examples as pairs of  $(x, y)$  values, such as (house size, house price).
- The model learns to predict  $y$  from  $x$  by adjusting its parameters ( $w$  and  $b$ ).
- After training, we can predict  $y$  values for new  $x$  inputs the model hasn't seen before.





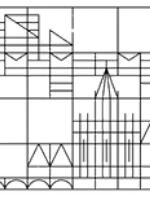
# Loss Function

## Mean Absolute Error



The Loss Function provides a quantitative measure of how well model predictions align with actual targets.

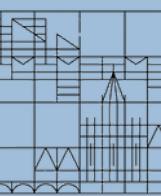
- They are a clear optimization target.
- We aim to find the combination of model parameters that minimizes the chosen loss function.
- Different problems call for different types of loss functions.
  - Mean Squared Error (MSE) or Mean Absolute Error (MAE) are typical choices for regression
  - Cross Entropy Loss is the typical choice for classification



# Notation

I will try to be consistent with the notation in the slides

- $x$  will always be the input
- $y$  the output
- bold letters denote vectors
  - for instance  $\mathbf{x}$  is an input with multiple dimensions
- bold capital letters denote matrices
  - for instance  $\mathbf{W}$  can be the weight matrix of a Neural Network
  - $\mathbf{Wx}$  is the product between the matrix  $\mathbf{W}$  and the vector  $\mathbf{x}$
- the parameters of the models are denoted by Greek letters or with a  $\mathbf{W}$
- the loss function is denoted with an  $L$ 
  - it is a function of the model parameters  $L[\mathbf{W}]$



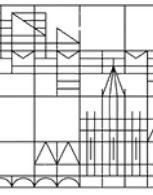
## *Math Recap*

# Matrix Multiplication

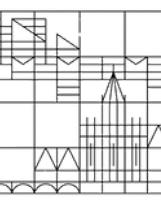
Matrix Multiplication consists in row by column multiplications:

- the elements  $(i, j)$  of the product matrix is obtained starting from the row  $i$  of the first matrix and the column  $j$  of the second matrix
- in general matrix multiplication is non commutative  $\mathbf{A} \times \mathbf{B} \neq \mathbf{B} \times \mathbf{A}$
- the number of columns of the first matrix must be equal to the number of rows of the second matrix

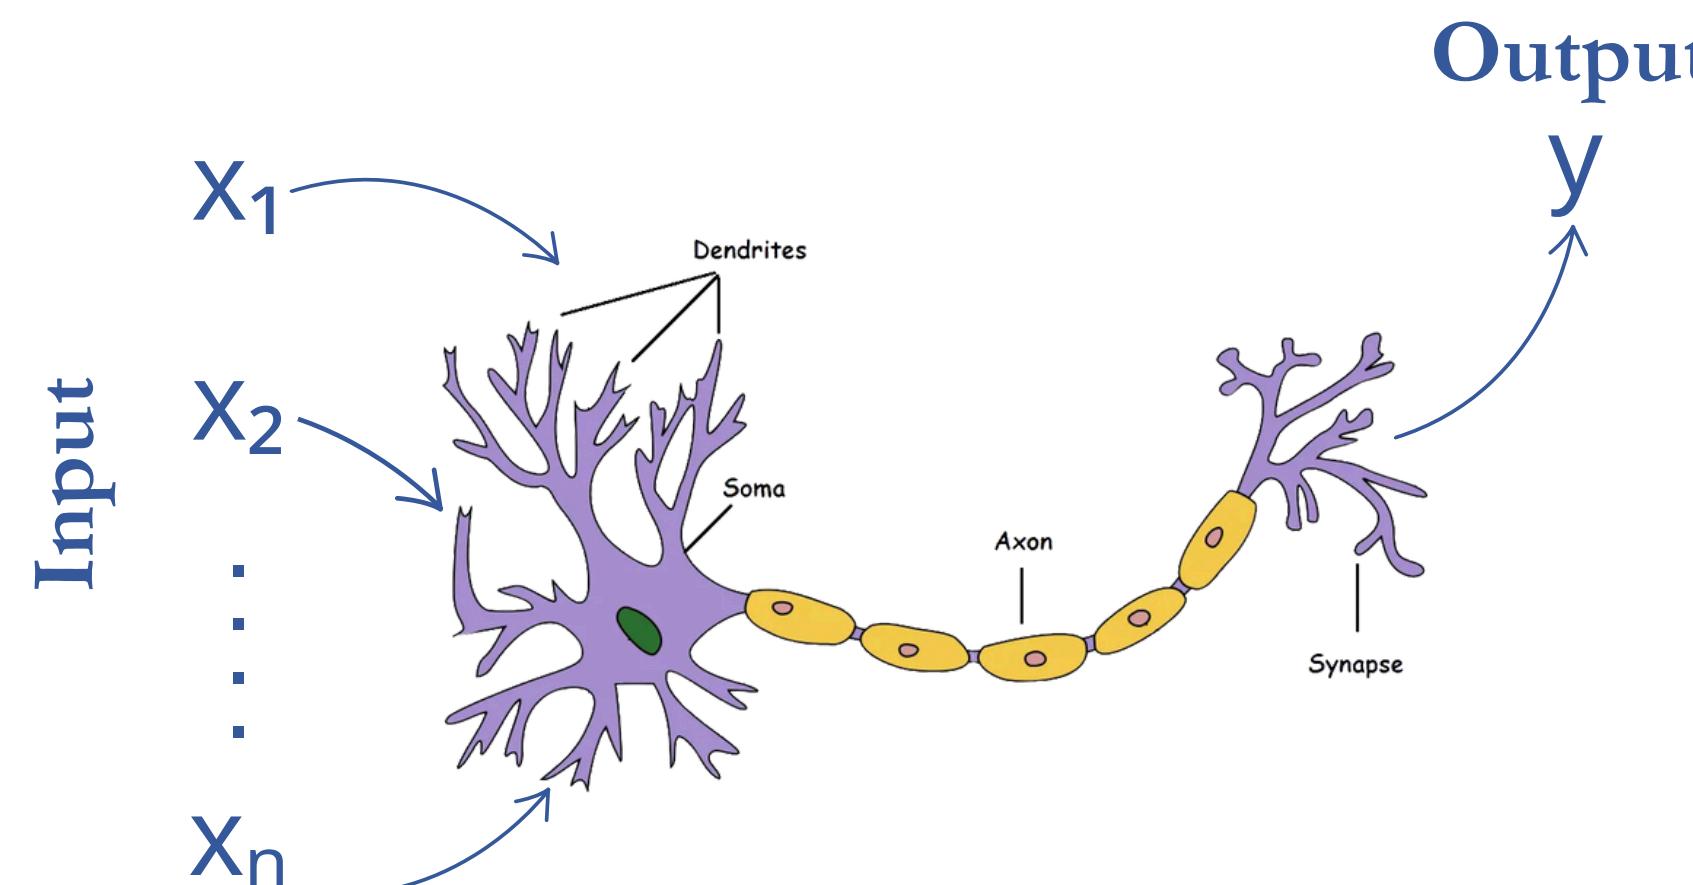
$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 10 & 11 \\ 20 & 21 \\ 30 & 31 \end{bmatrix} = \begin{bmatrix} 1 \times 10 + 2 \times 20 + 3 \times 30 & 1 \times 11 + 2 \times 21 + 3 \times 31 \\ 4 \times 10 + 5 \times 20 + 6 \times 30 & 4 \times 11 + 5 \times 21 + 6 \times 31 \end{bmatrix} = \begin{bmatrix} 10+40+90 & 11+42+93 \\ 40+100+180 & 44+105+186 \end{bmatrix} = \begin{bmatrix} 140 & 146 \\ 320 & 335 \end{bmatrix}$$



# The Perceptron



# Biological Neurons



Neural networks draw inspiration from the brain's biological structure.

- **Input:** Neurons receive signals from other neurons through dendrites.
- **Processing:** Incoming signals are combined and processed in the cell body.
- **Activation Function:** If the processed signal is strong enough, it triggers firing.
- **Output:** When activated, the neuron sends a signal through its axon.



# The First Neural Network

The Perceptron was the first artificial neural network model.

- Created in the 1950s by Frank Rosenblatt, a pioneer in artificial intelligence
- Simulates how neurons in the human brain process information

The Perceptron follows a simple operational principle:

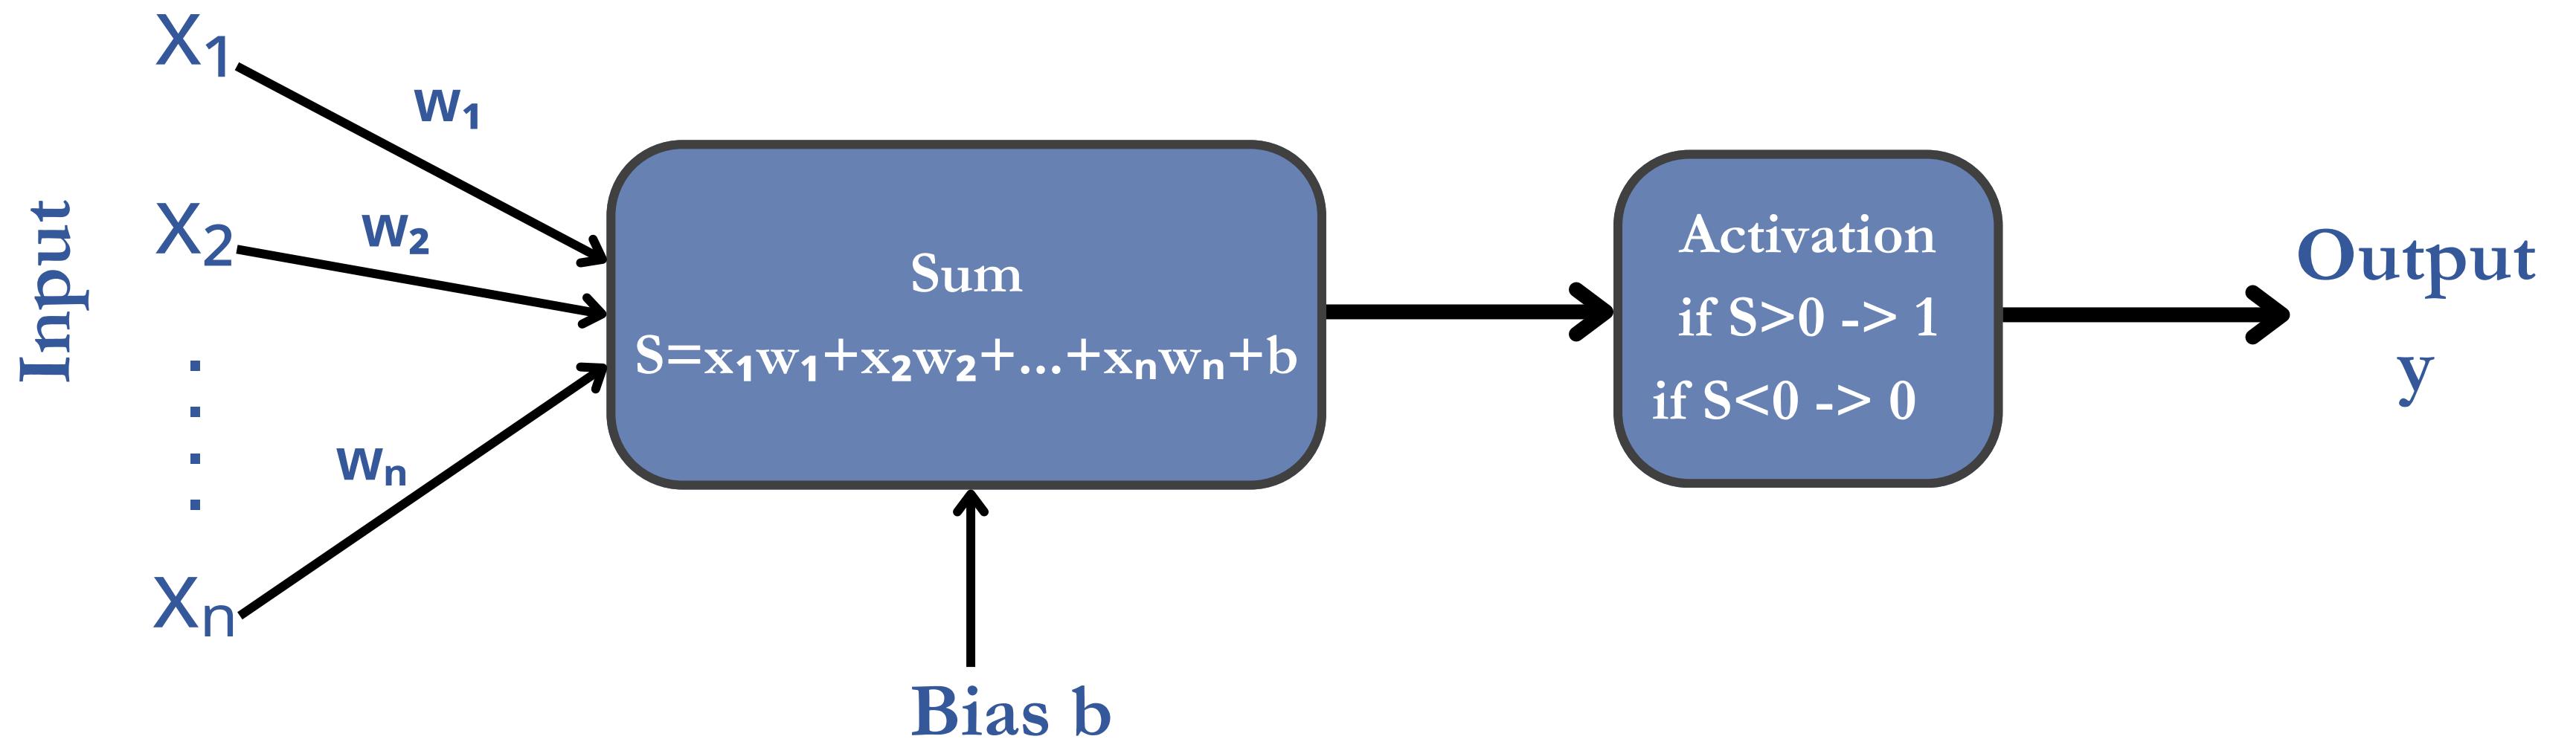
- It receives multiple inputs, each with an associated weight
- These inputs are combined and produce an output if they exceed a threshold

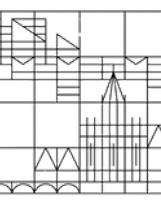




# Structure of the Perceptron

The Perceptron combines weighted inputs and activation:





# Mathematical Notation

Mathematically we can represent the perceptron using a vector product

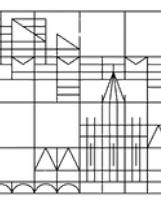
$$y = a[\mathbf{w}\mathbf{x} + b]$$

Here

- $\mathbf{w} = (w_1, w_2, \dots, w_n)$  is the vector containing the  $n$  weights
- $\mathbf{x} = (x_1, x_2, \dots, x_n)$  is the vector containing the  $n$ -dimensional input
- $a$  is the activation function. In our case
  - $a(x) = 1$  if  $x > 0$
  - $a(x) = 0$  if  $x < 0$
- $b$  is the bias

With  $\mathbf{w}\mathbf{x}$  we denote the scalar product of the two vectors

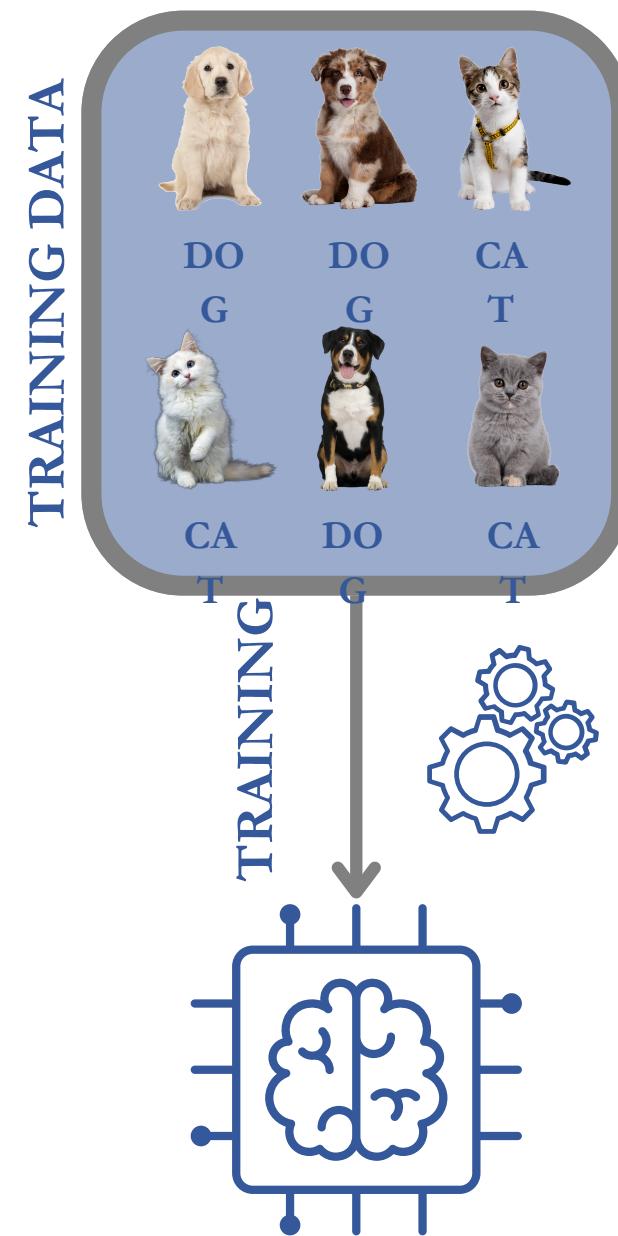
$$\mathbf{w}\mathbf{x} = x_1w_1 + x_2w_2 + \dots + x_nw_n$$



# Supervised Classification

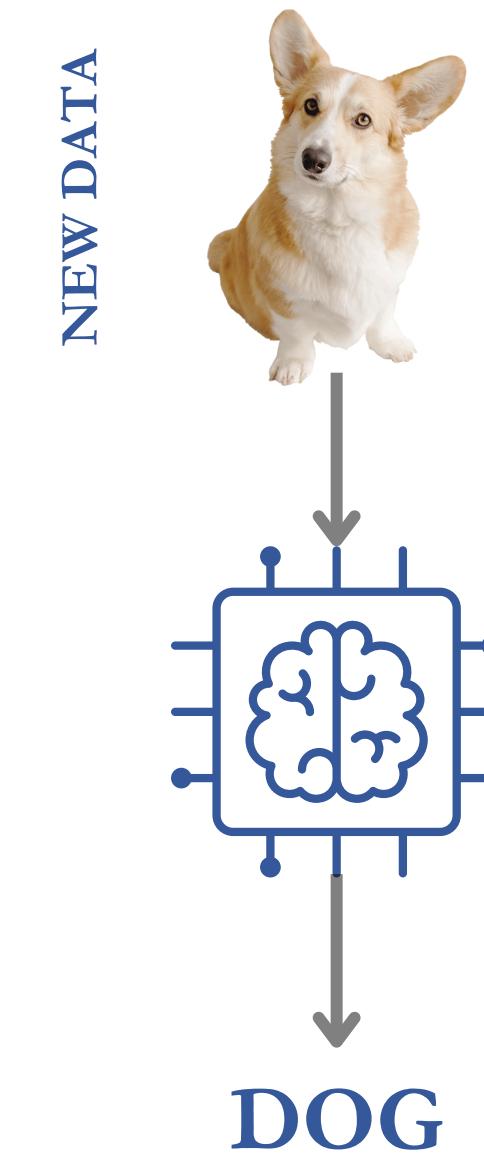
The computer learns from classified examples to predict categories for new data.

## LEARNING

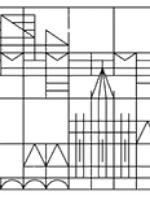


The model analyzes pre-classified examples and learns how to distinguish between categories

## CLASSIFICATION



The model applies what it learned to categorizes new, previously unseen data

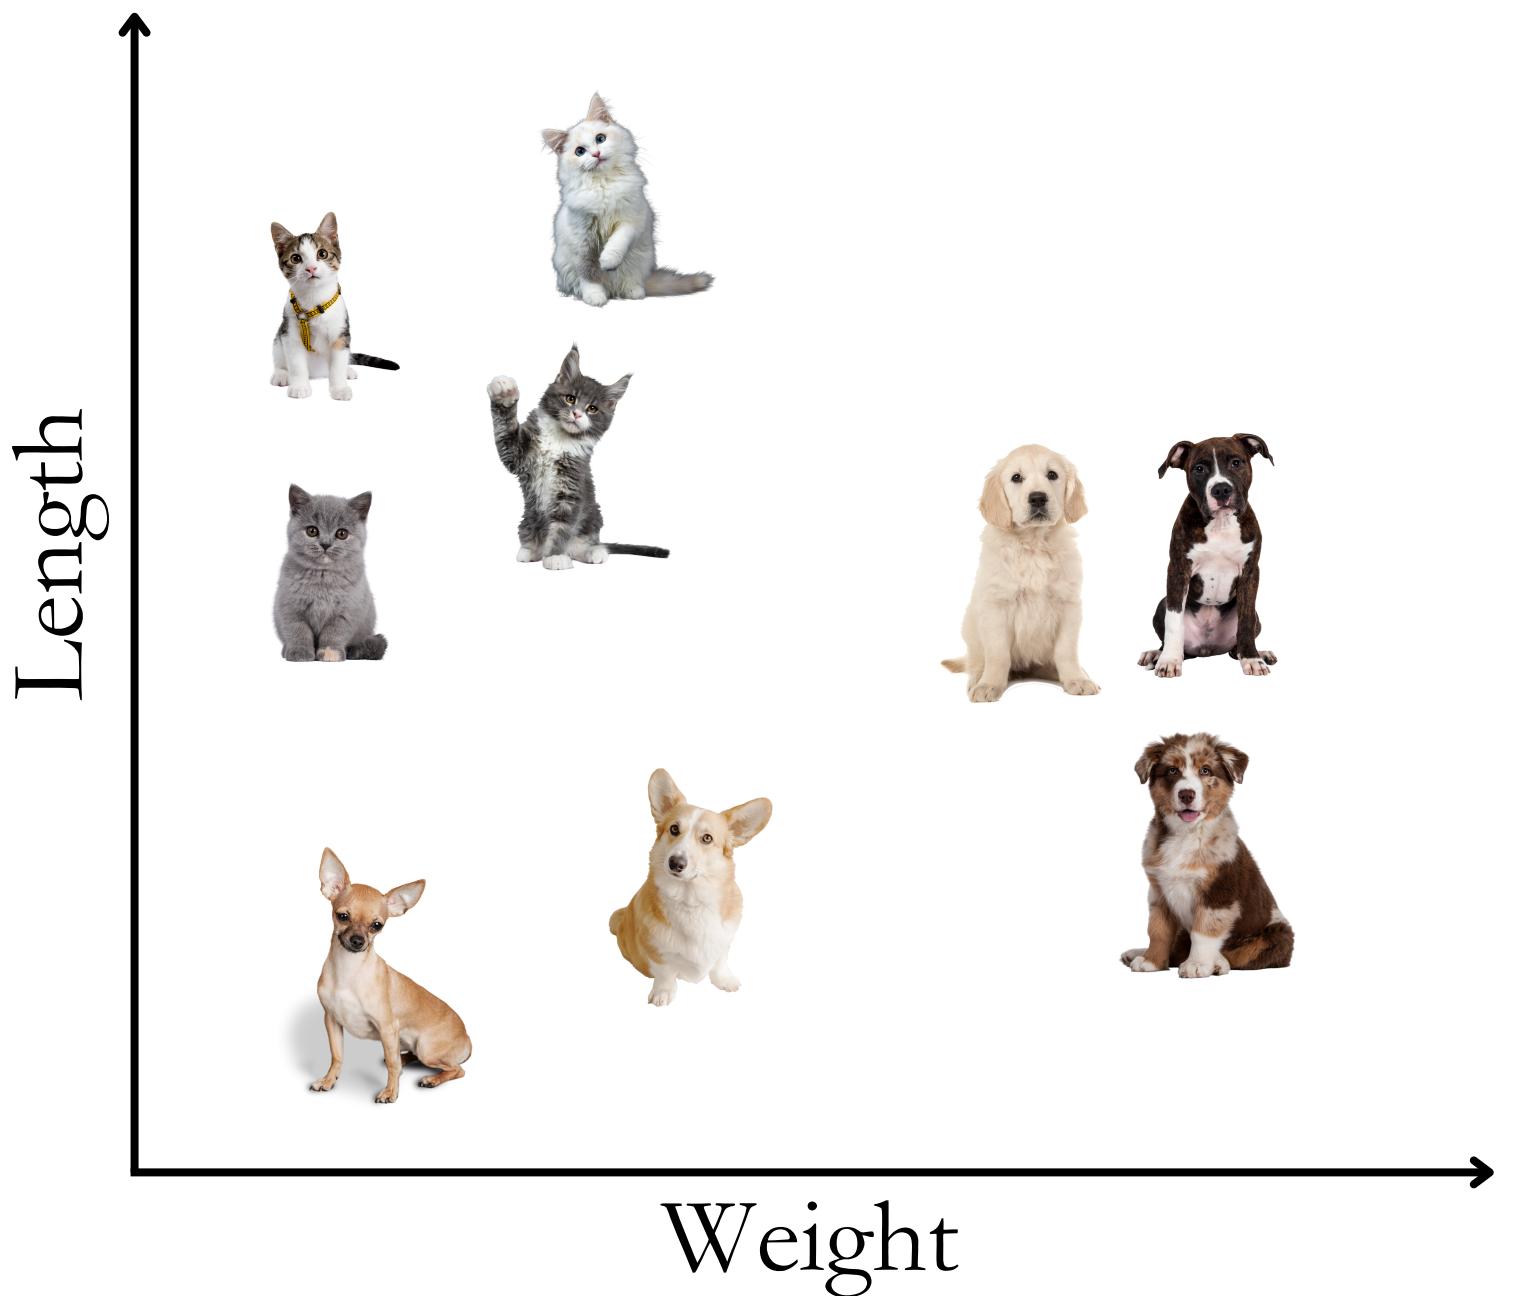


# Classifying Dogs and Cats

Let's consider a simplified version of the classification problem

- Dogs and cats can be characterized by various features
  - We focus on just two: weight and length
- When can plot different animals on a graph using weight and length
- Each animal is represented by a point

The Perceptron's job is to determine which group (dog or cat) a new animal belongs to

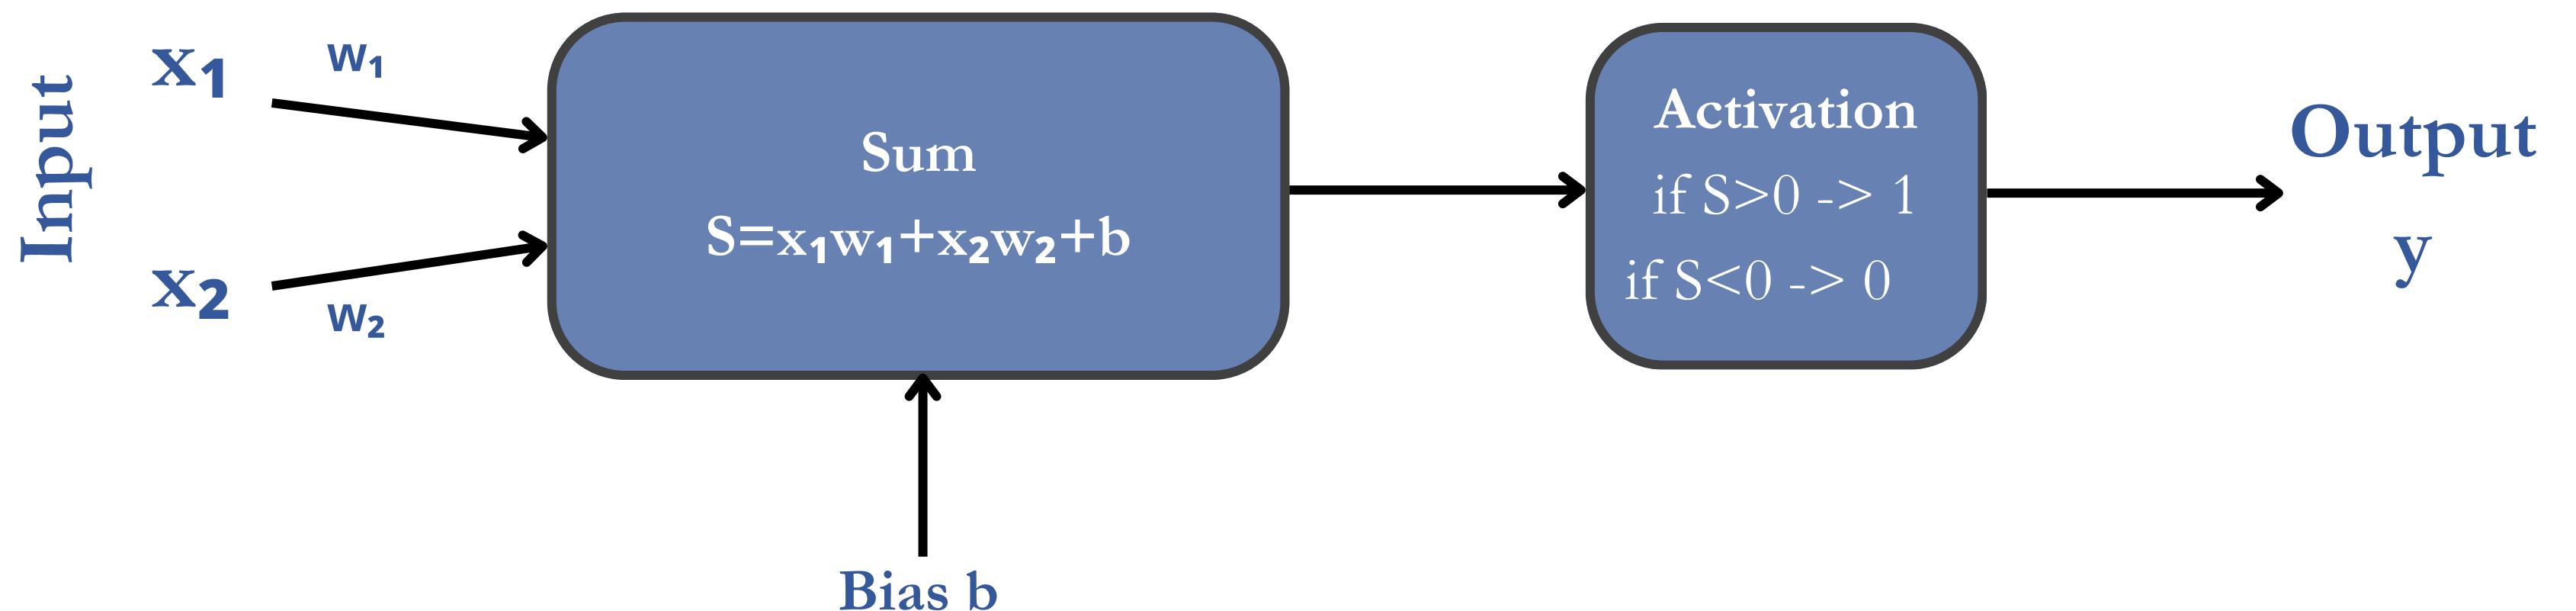




# Perceptron with 2 Inputs

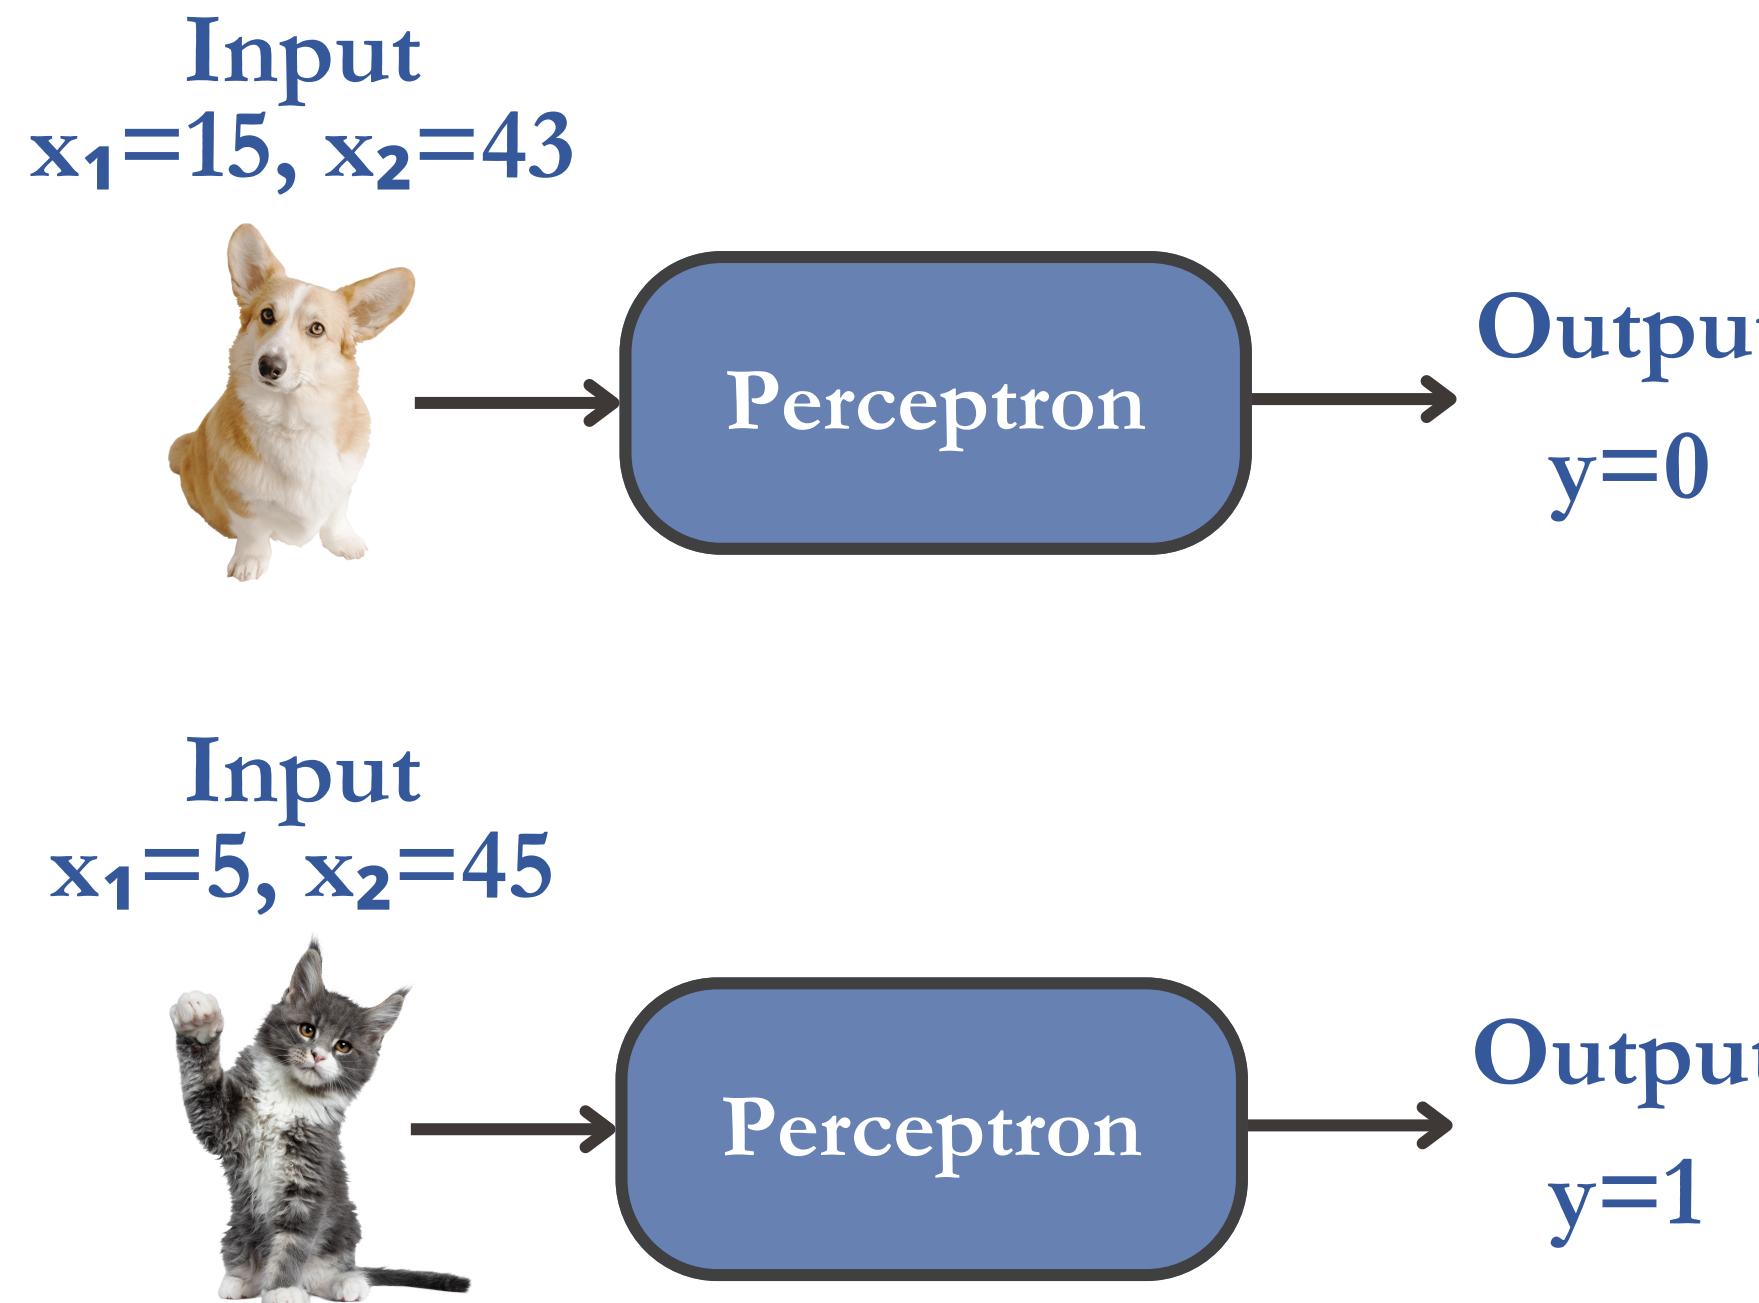
In this simple case we have just 2 inputs  $\mathbf{x}=(x_1, x_2)$  and 3 parameters

- $x_1$  is the weight of the animal,  $x_2$  is length
- $w_1$  and  $w_2$  are the weights
- the bias is denoted by  $b$



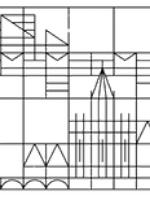


# Automatic Classification



The Perceptron distinguishes between two categories:

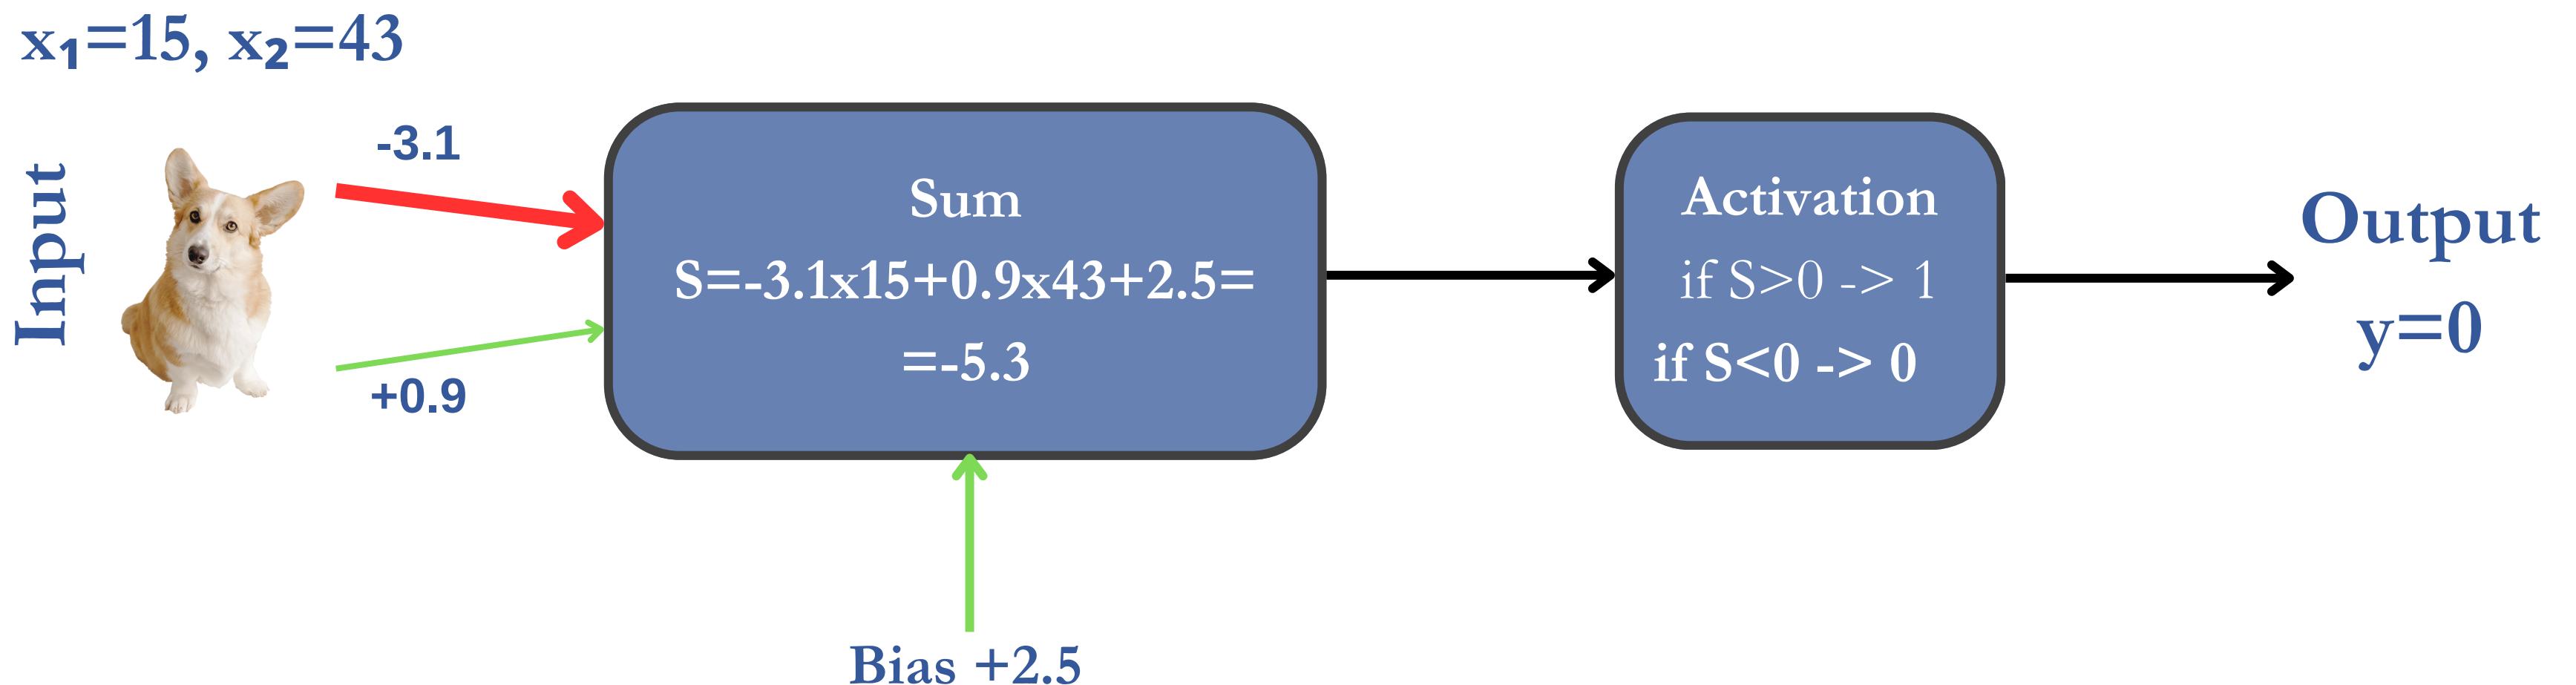
- **Input:** Animal characteristics (e.g., weight and length)
- **Processing:** The Perceptron applies weights and calculates the output
- **Output:** Classification result
  - $0 = \text{dog}$
  - $1 = \text{cat}$

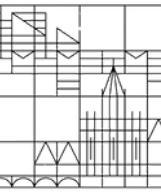


# Perceptron with 2 Inputs

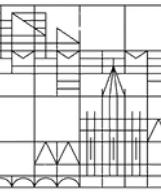
Let's see a practical example plugging some numbers in the perceptron

- we use as parameters  $w_1=3.1$ ,  $w_2=-0.9$  and  $b=-2.5$
- the input is  $\mathbf{x}=(15, 43)$





# Limits of the Perceptron



# Understanding the Perceptron

The Perceptron creates a linear decision boundary to separate categories.

**Cat:**  $w_1x + w_2y + b > 0$

**Dog:**  $w_1x + w_2y + b < 0$

The decision boundary itself is defined by:

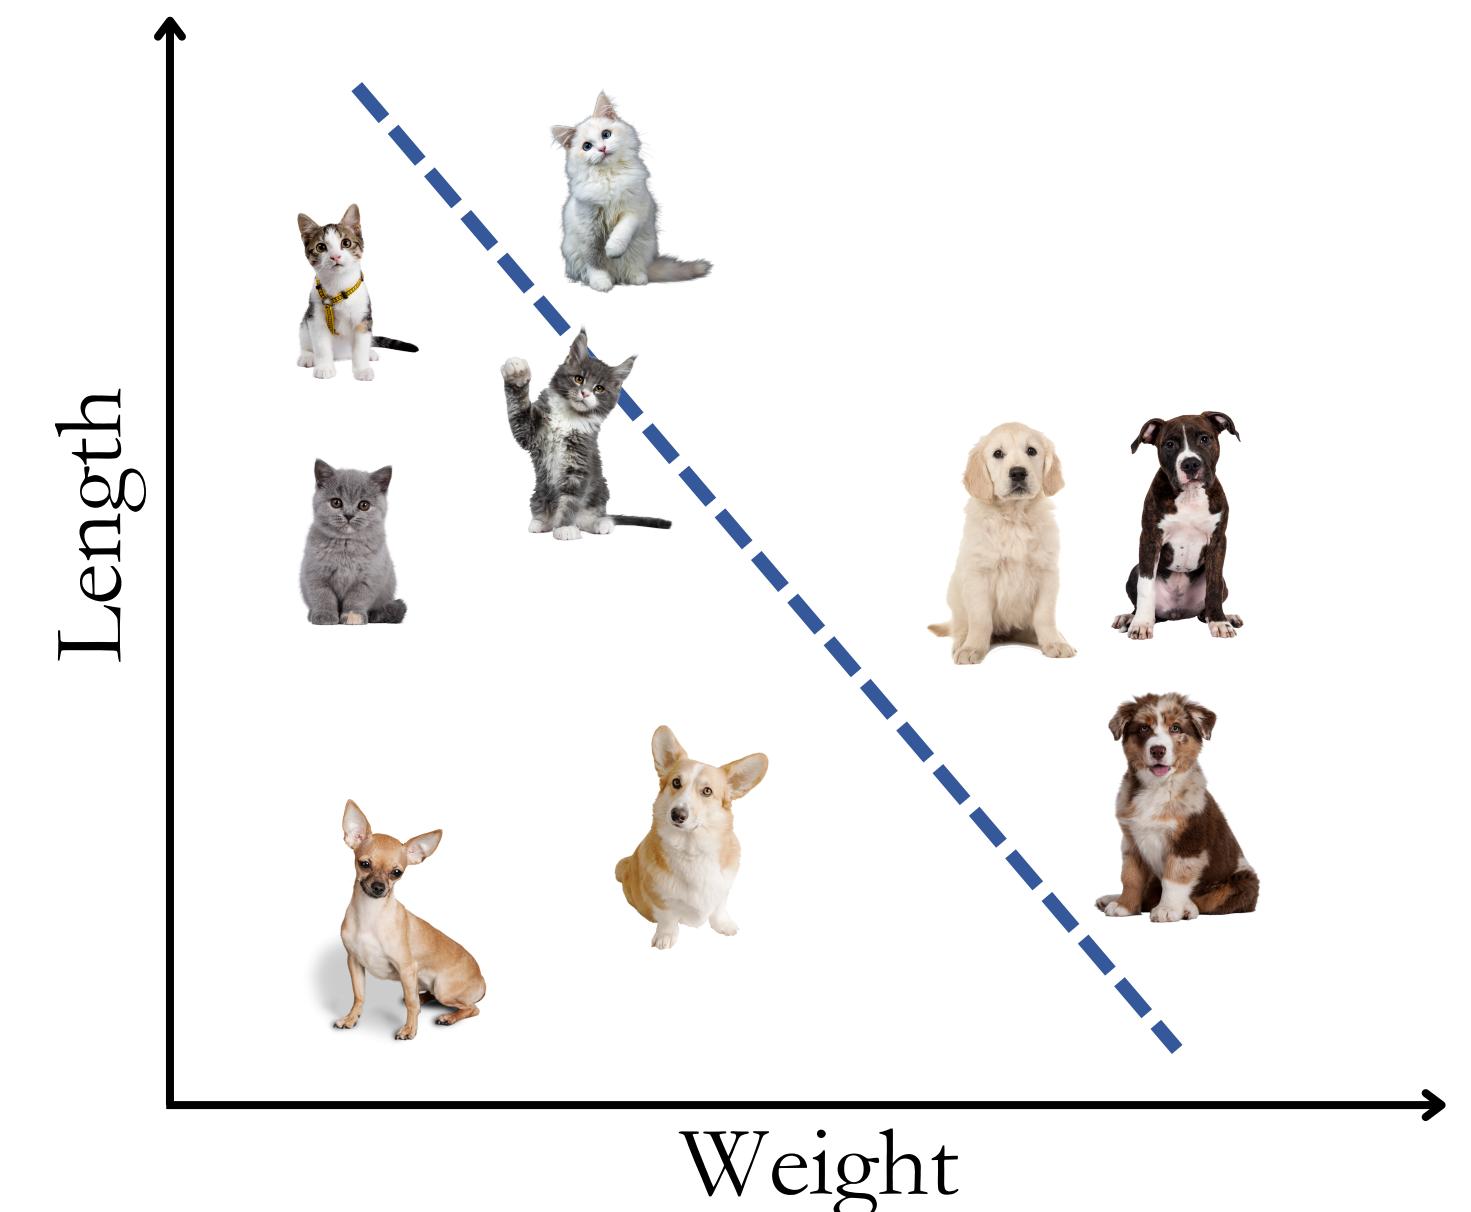
$$w_1x + w_2y + b = 0$$

Which can be rewritten as:

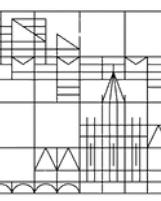
$$y = -(w_1/w_2)x - (b/w_2)$$

This demonstrates that the Perceptron:

- Draws a straight line in a 2D feature space
- Classifies points above the line as cats
- Classifies points below the line as dogs

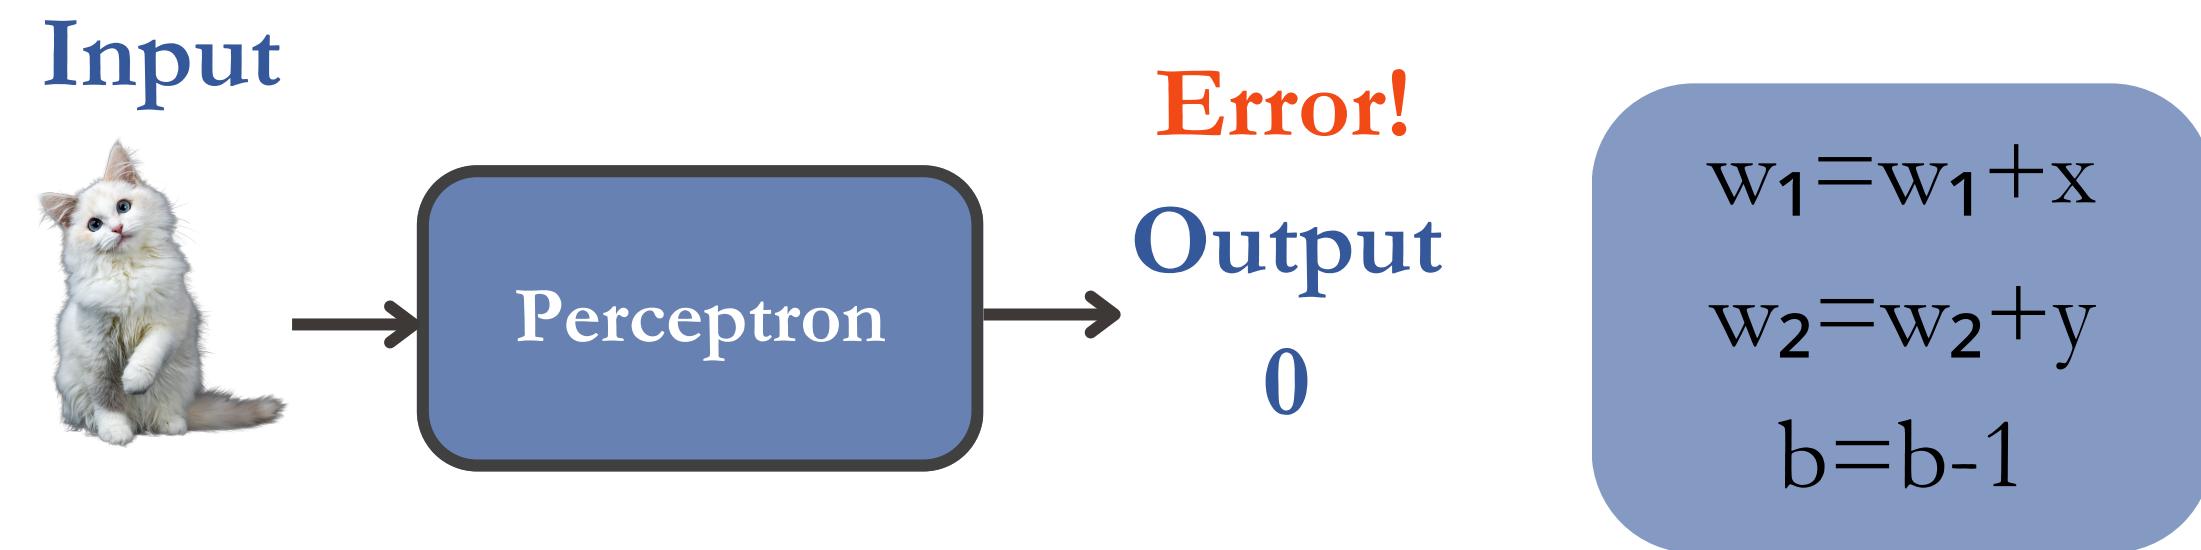
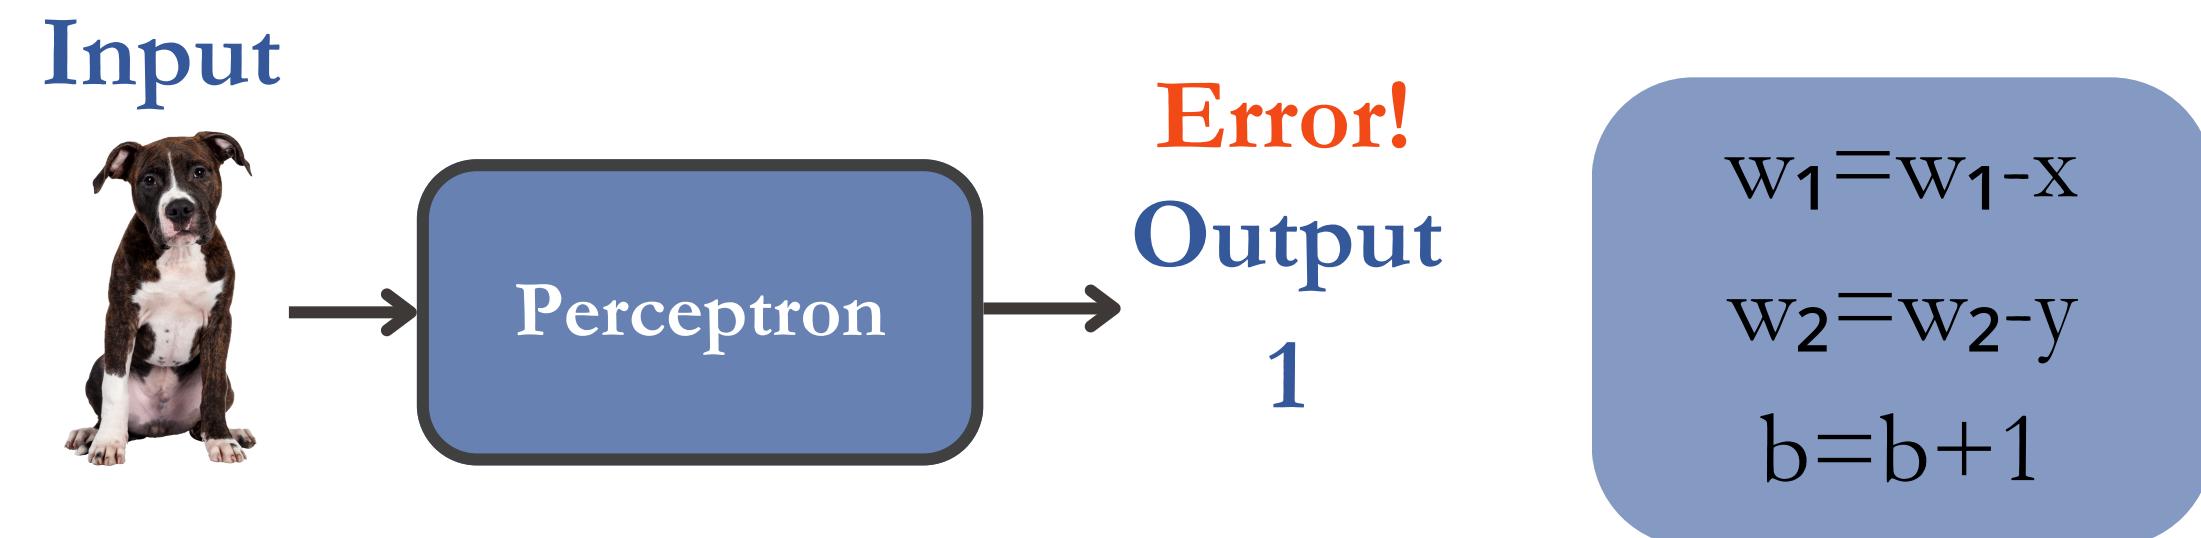


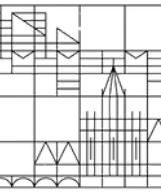
<https://giordano-demarzo.github.io/teaching/deep-learning-25/perceptron/>



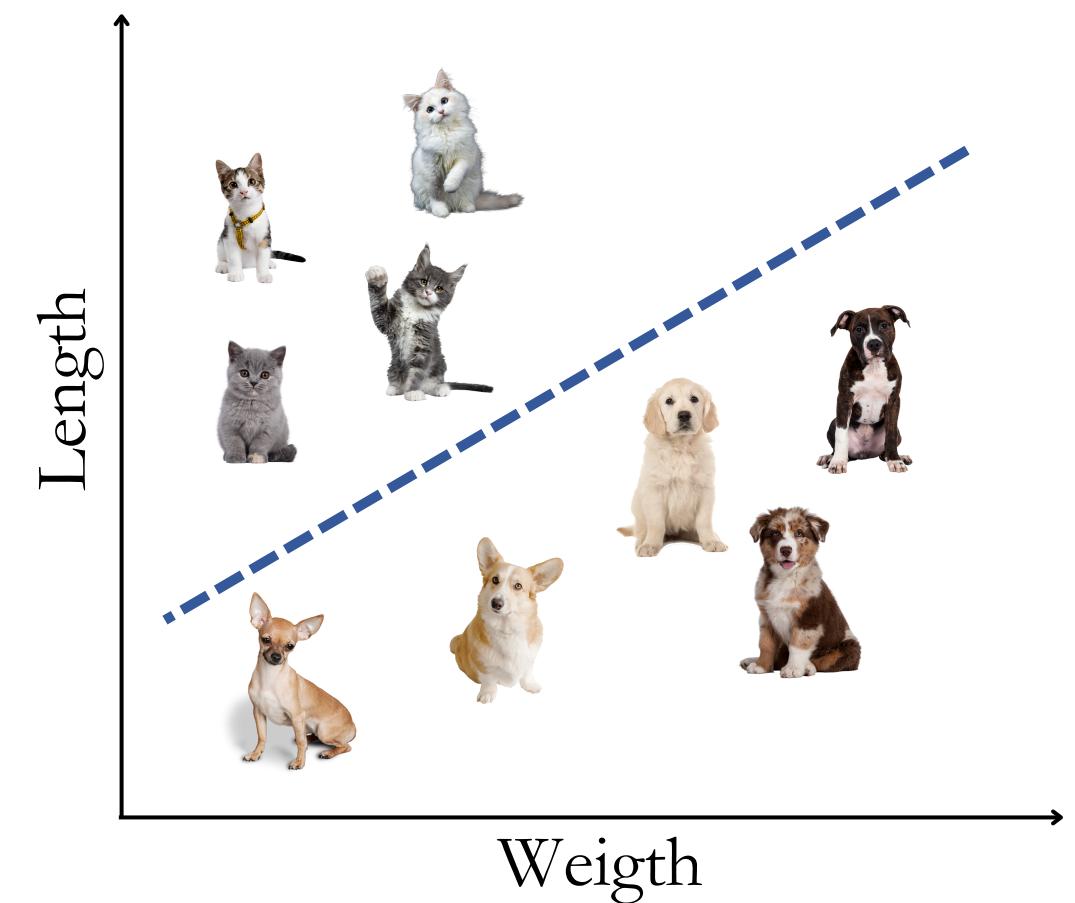
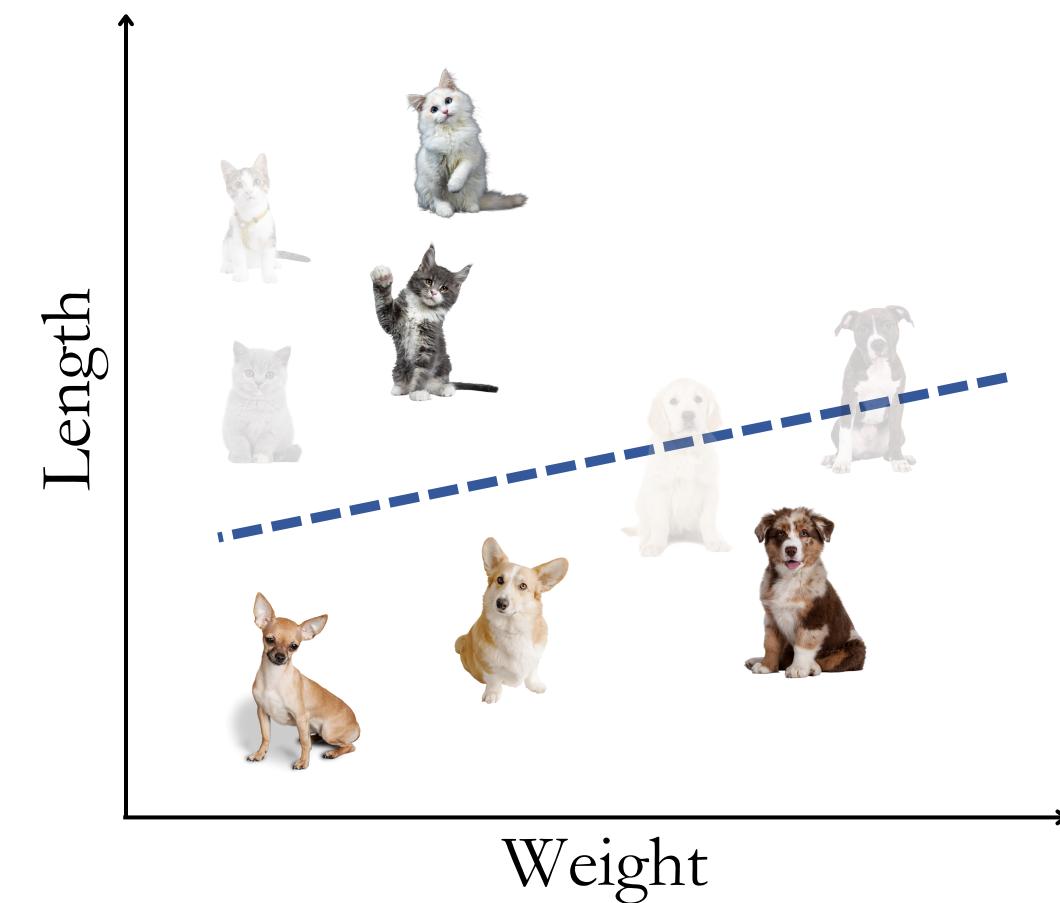
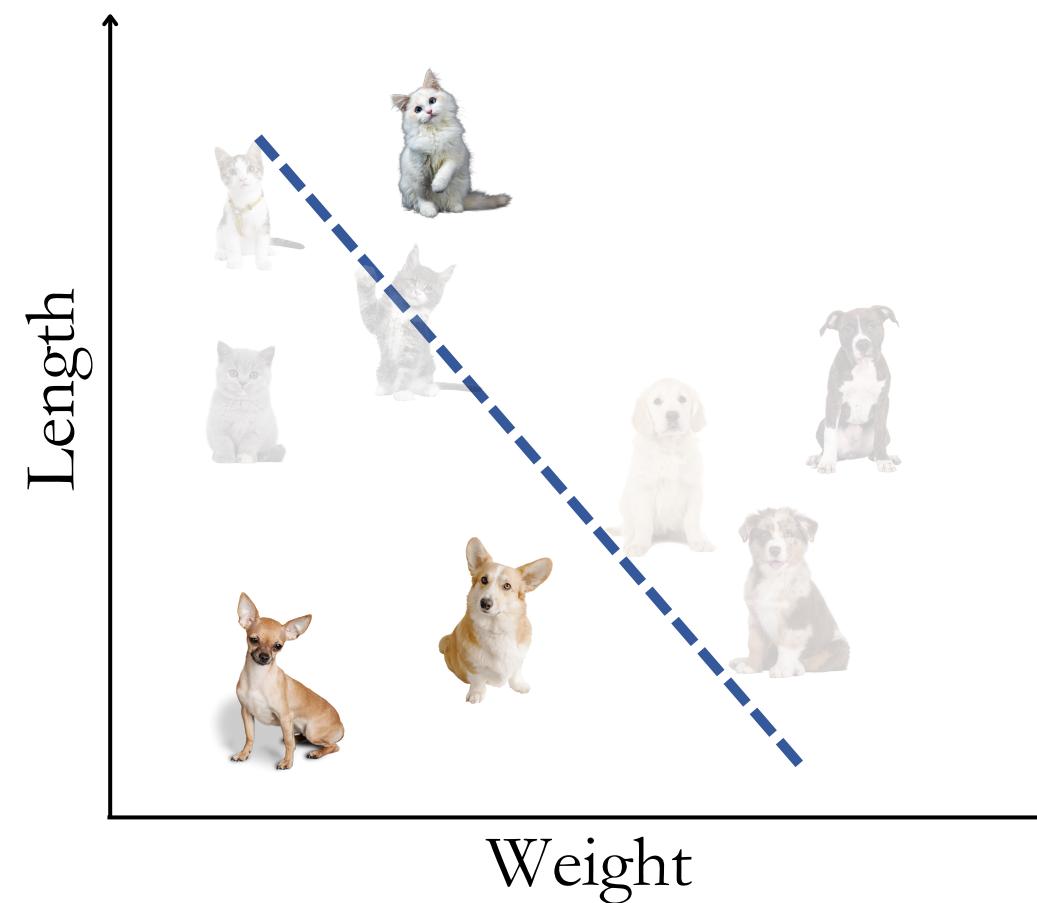
# Training Rule

During training the perceptron is shown labelled data and its weights are adjusted when it produces wrong classifications

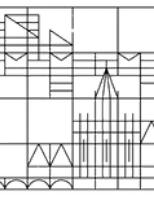




# Visualizing Training

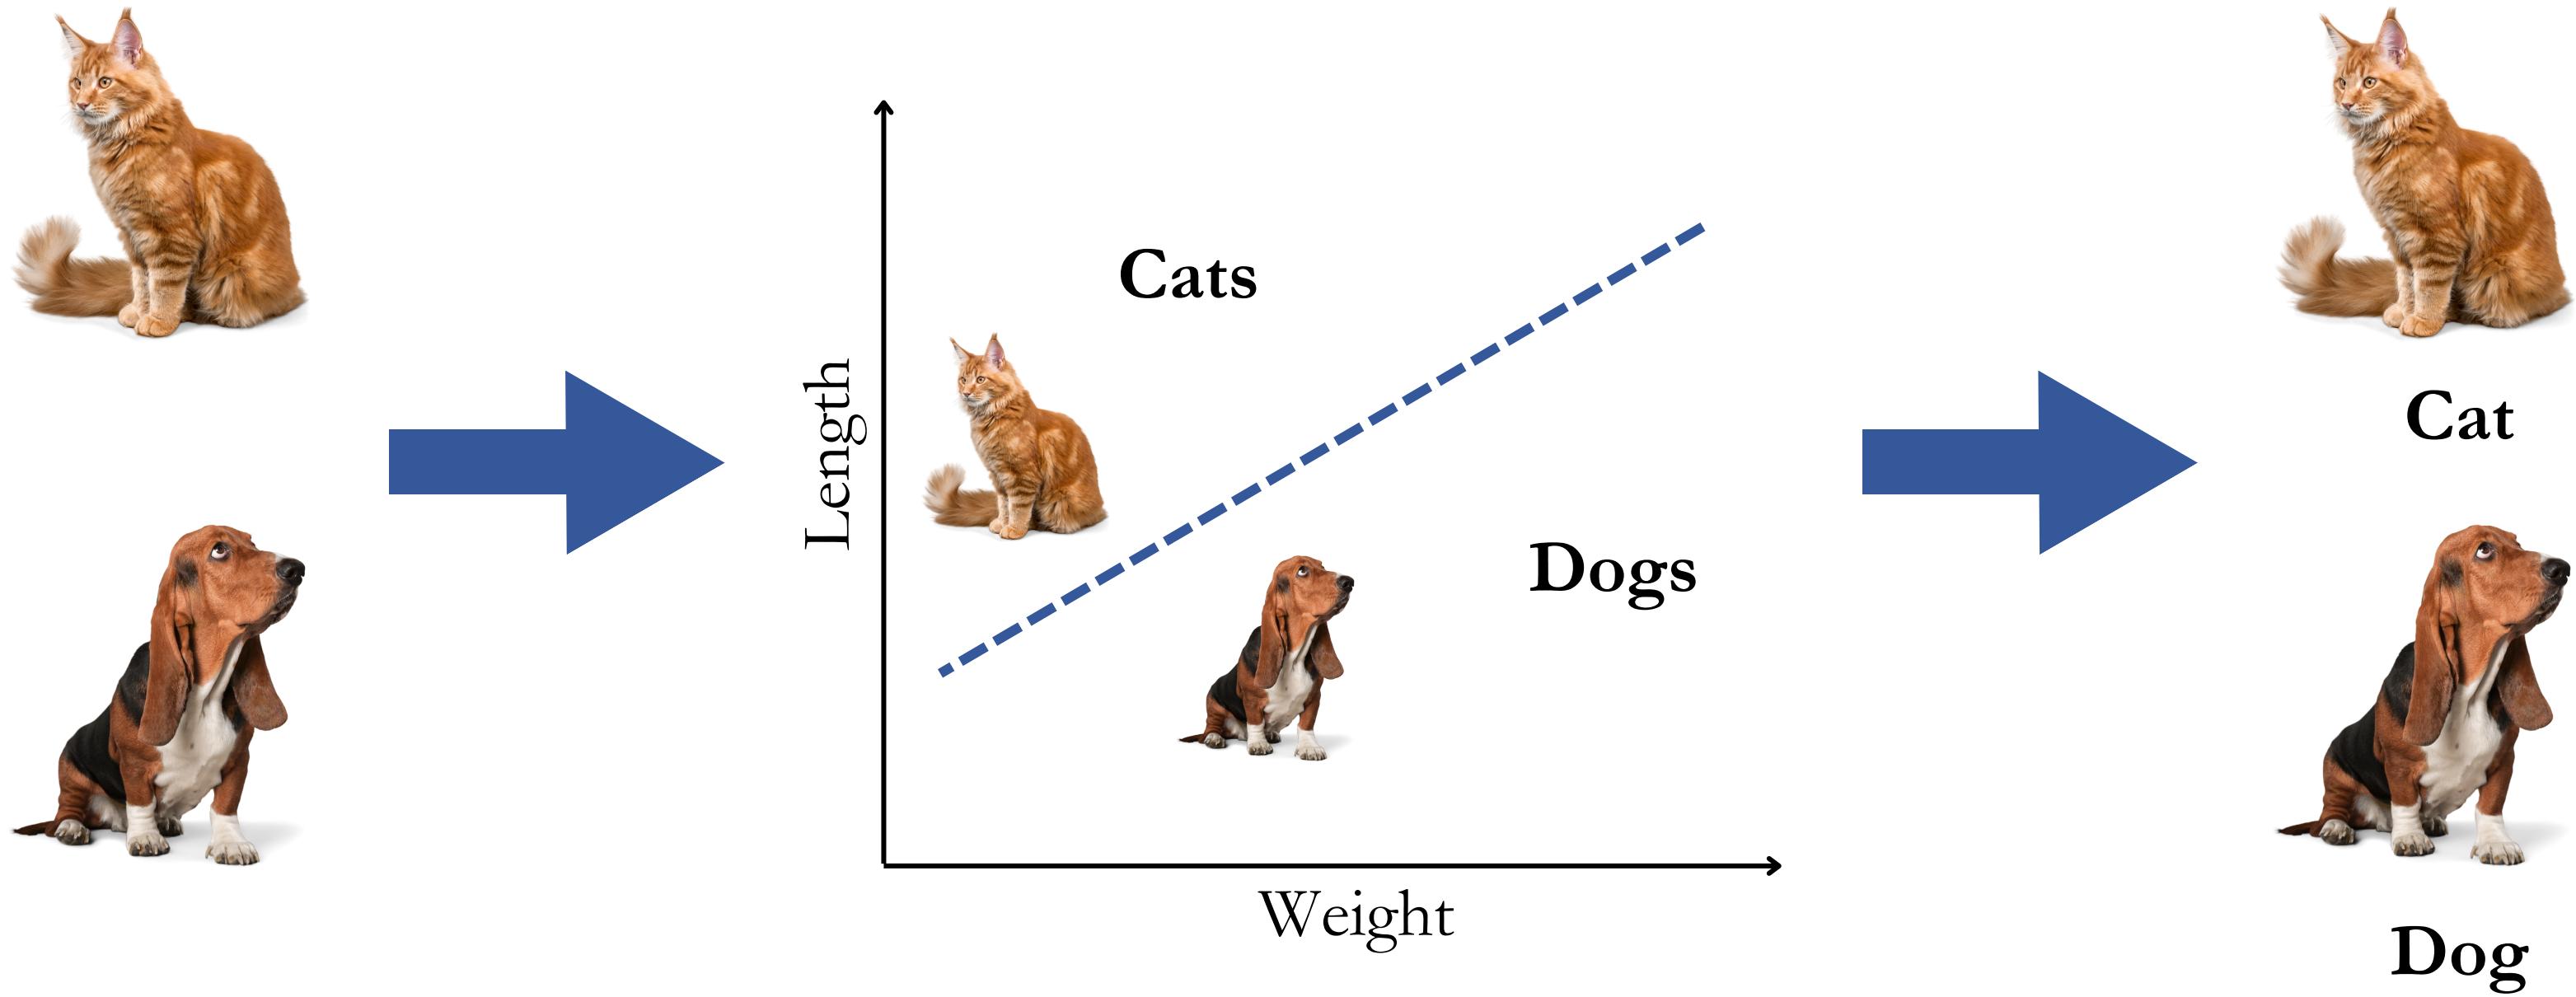


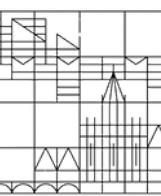
As training progresses, the decision boundary moves to better separate the classes.



# Classification

Once trained, the Perceptron can classify new animals from their weight and length

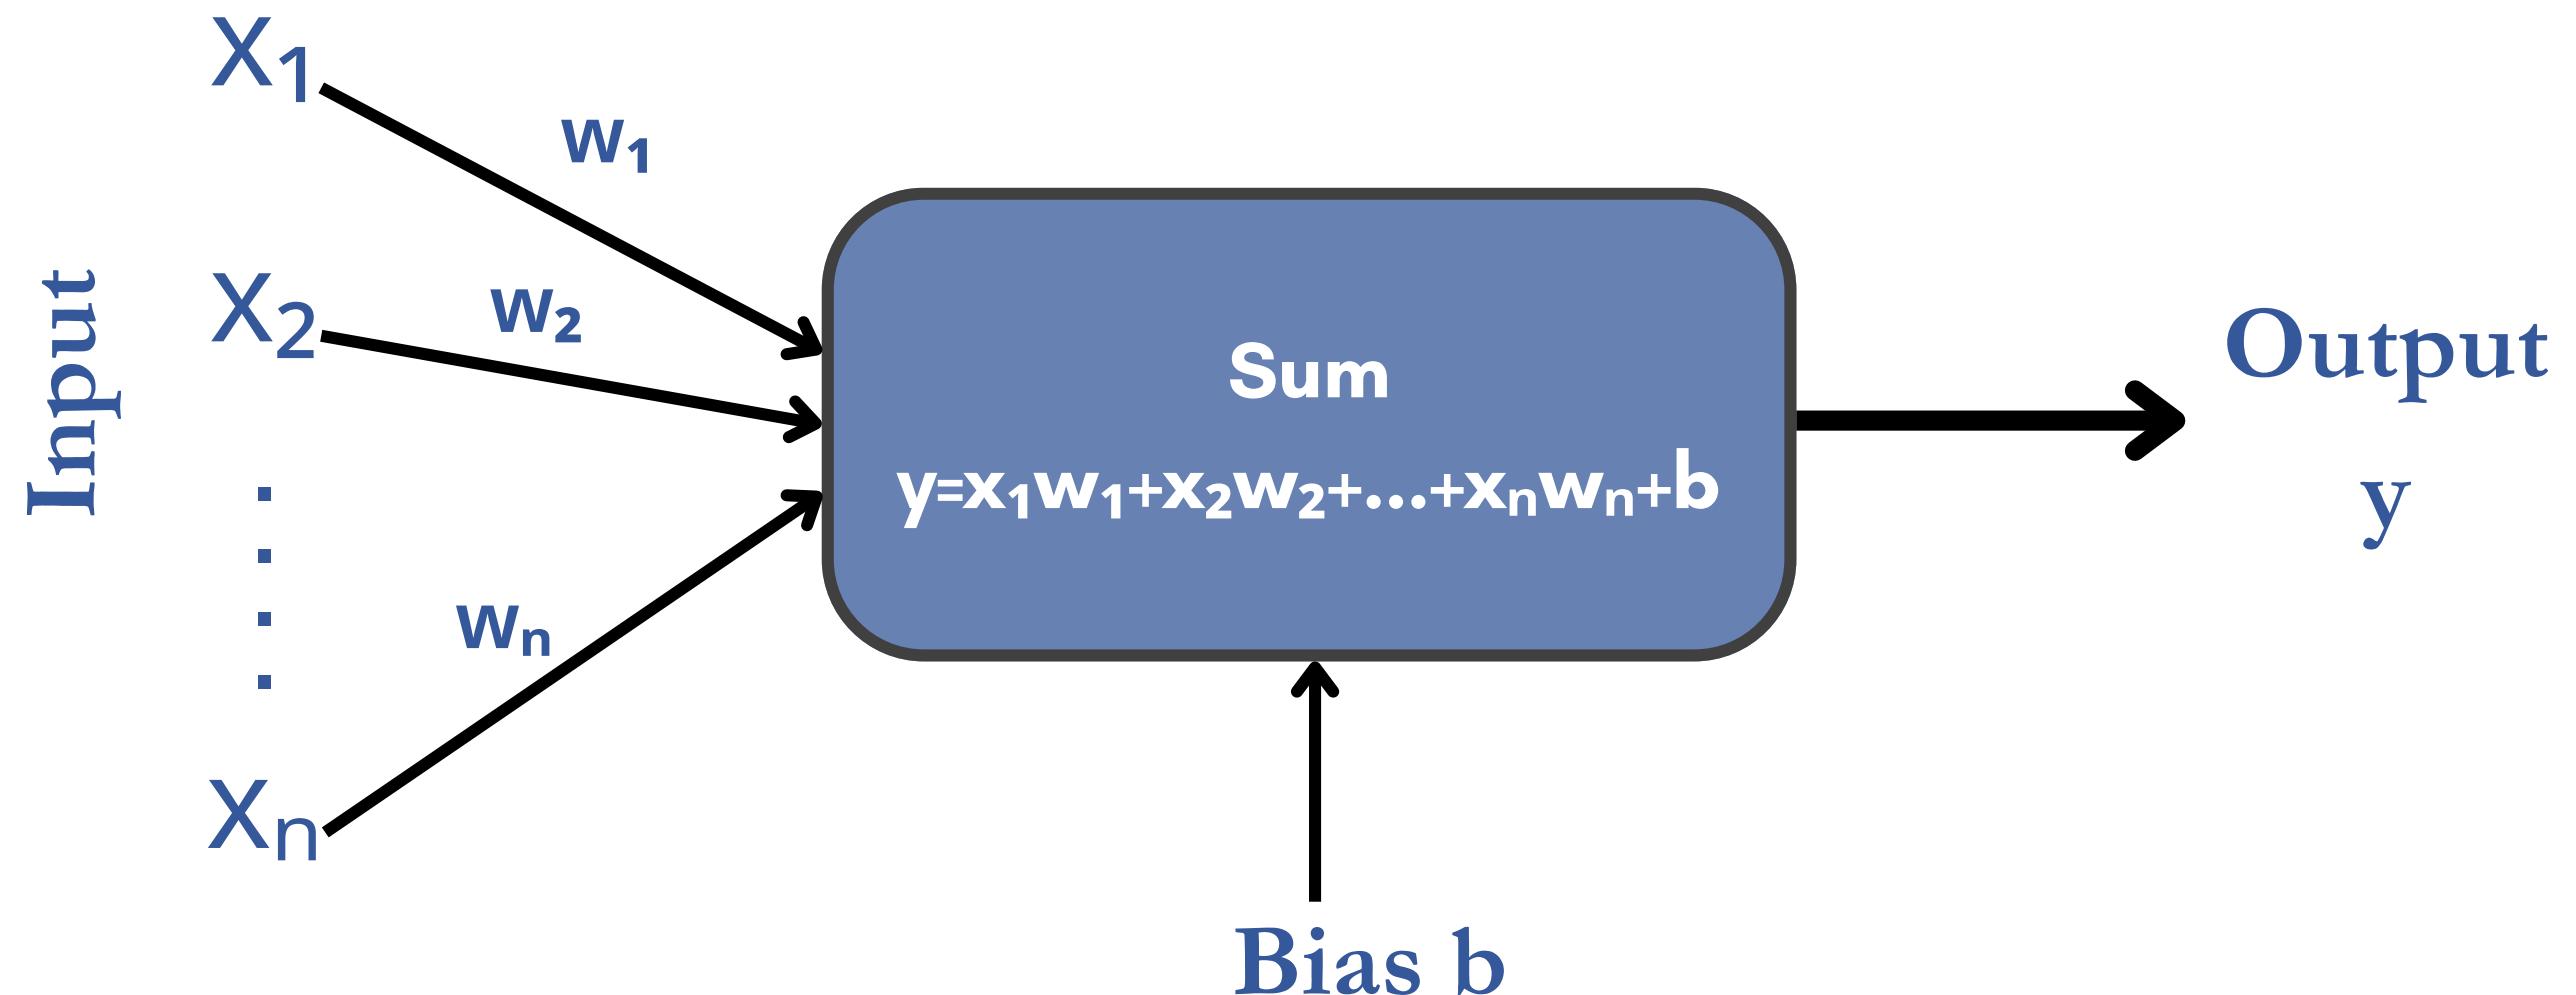


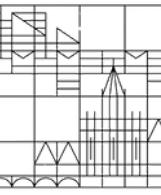


# Perceptron and Regression

The Perceptron can also perform regression

- Predicts a continuous numerical value instead of a category
- Skips the thresholding step in the activation function

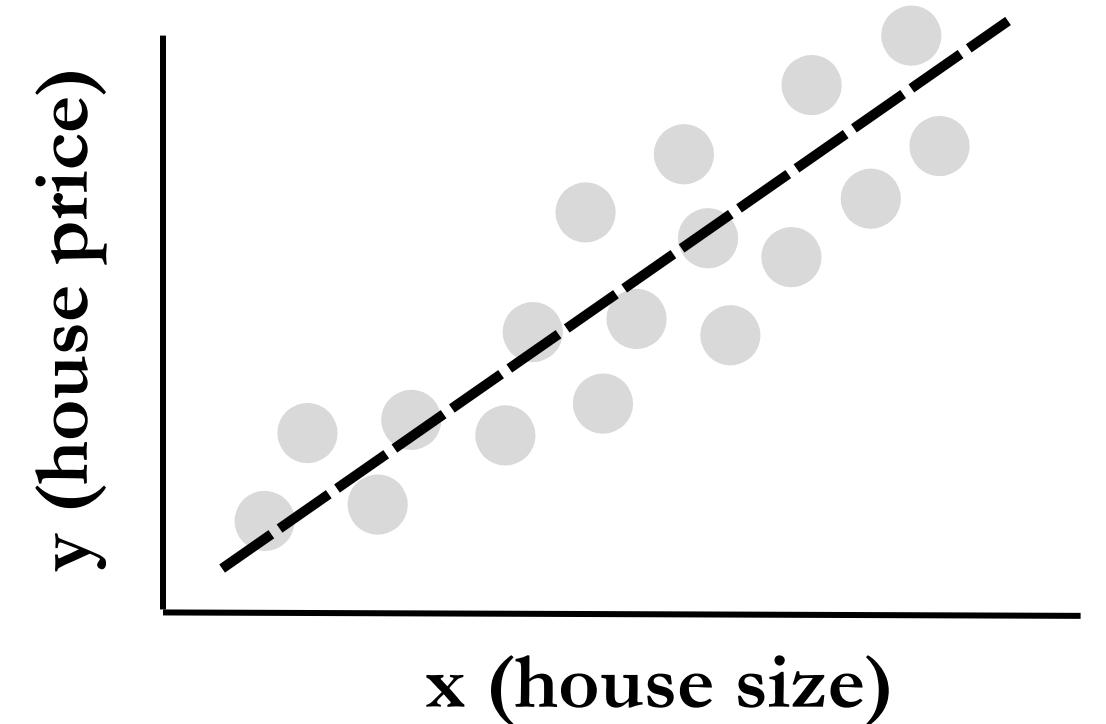
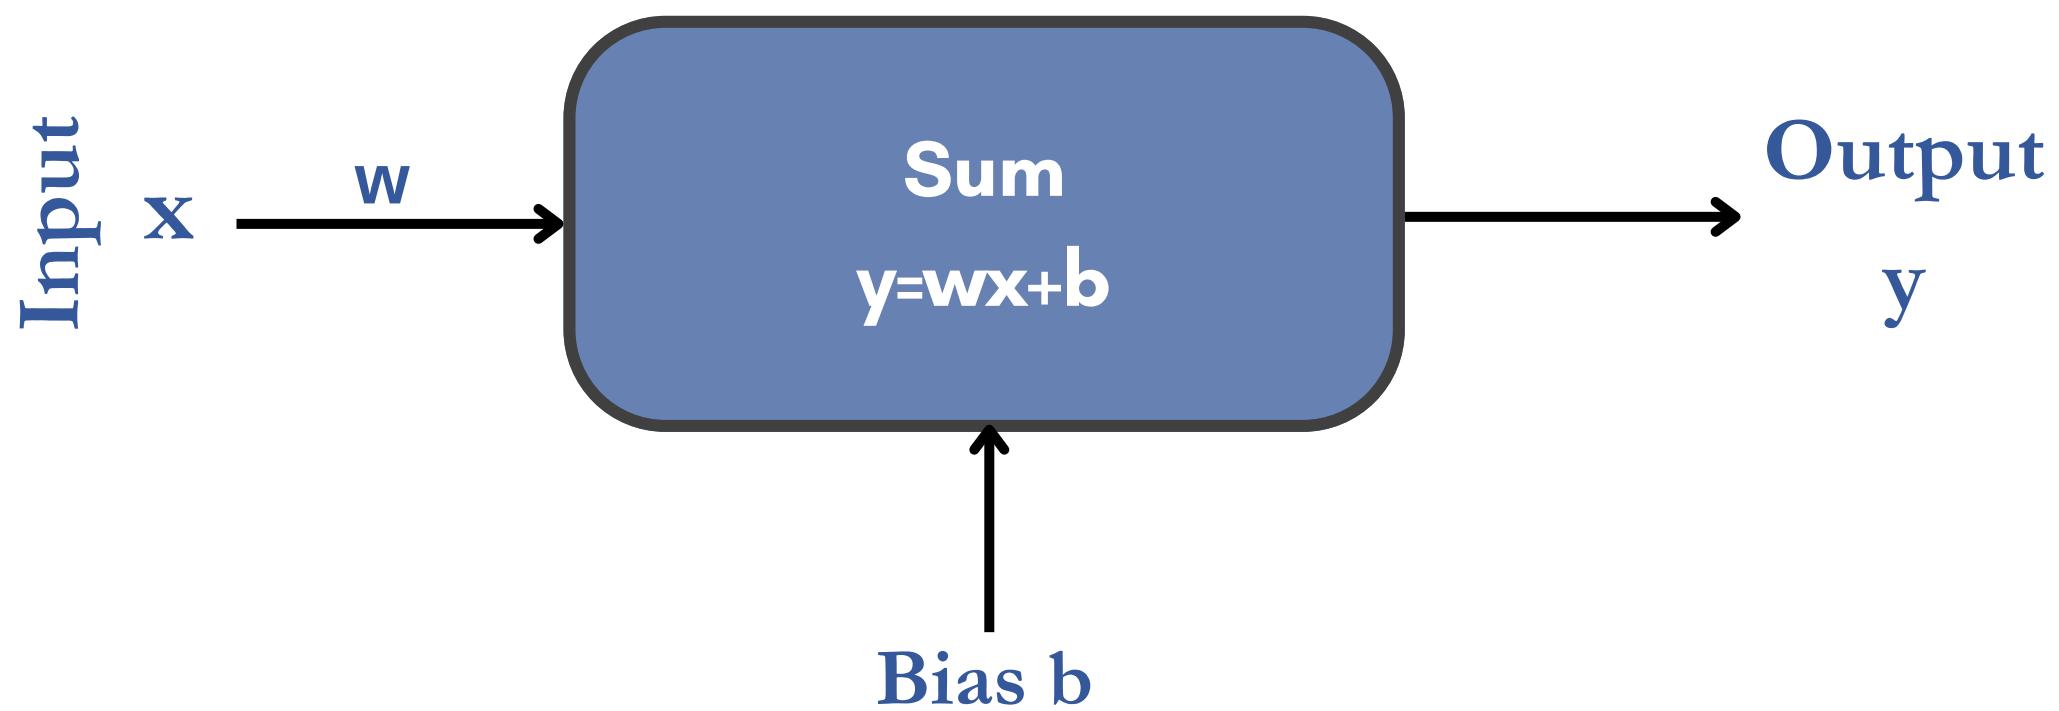


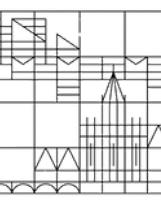


# Single Input Case

In the most simple case we have a single input

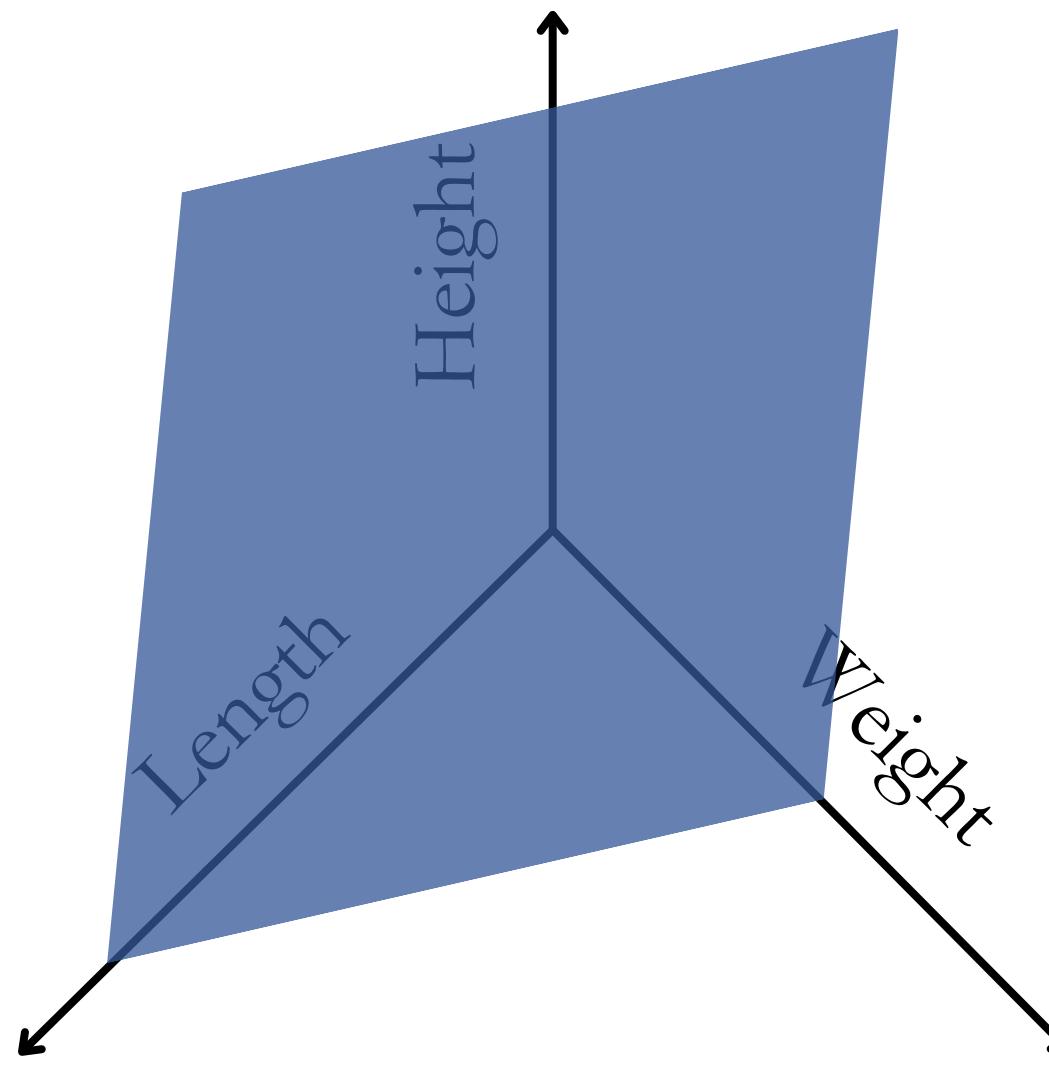
- the model output is  $y = wx + b$
- during training the model learns  $w$  and  $b$  to fit the data
- this is equivalent to a linear regression





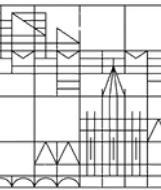
# Higher Dimensions

The Perceptron's principles extend to higher dimensions:



- 2 Dimensions: The Perceptron uses a line to separate categories (e.g., dogs and cats based on weight and length)
- 3 Dimensions: Adding another feature (e.g., height) creates a 3D space where the Perceptron uses a plane as separator

**What happens in higher dimensions?**



# Limits of the Perceptron

The Perceptron works well when data categories can be separated by a line (in 2D), a plane (in 3D), or a hyperplane (in higher dimensions).

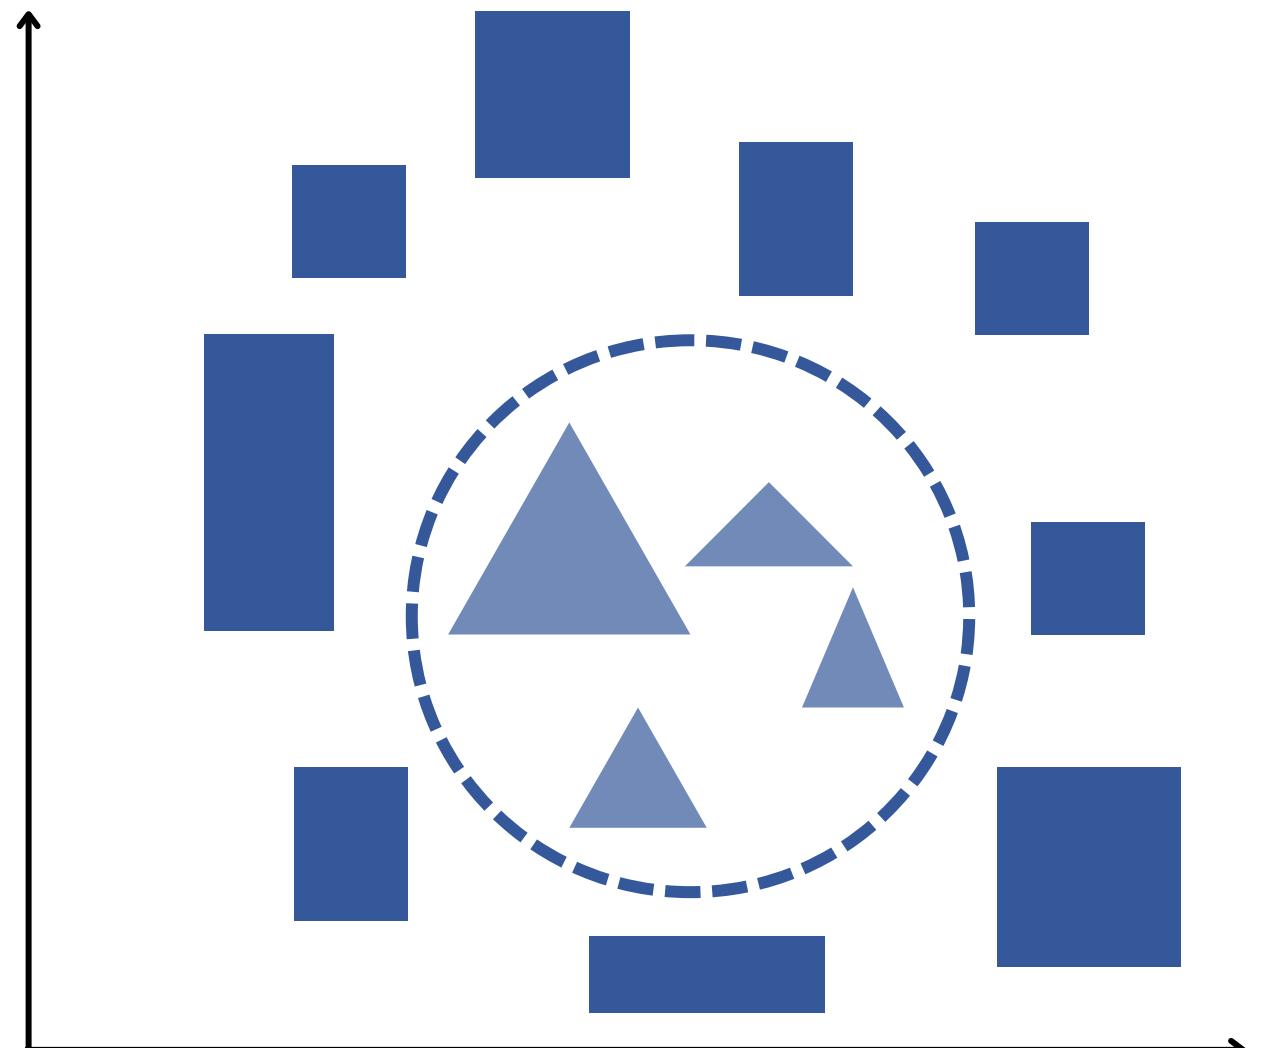
However, many real-world problems aren't linearly separable:

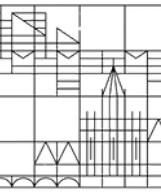
- If data forms patterns like circles or spirals
- If categories are intermingled in complex ways

In these cases, a single Perceptron is insufficient!

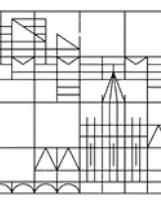
- Problems like XOR cannot be solved by a single Perceptron

This limitation led to the so called **AI Winter**





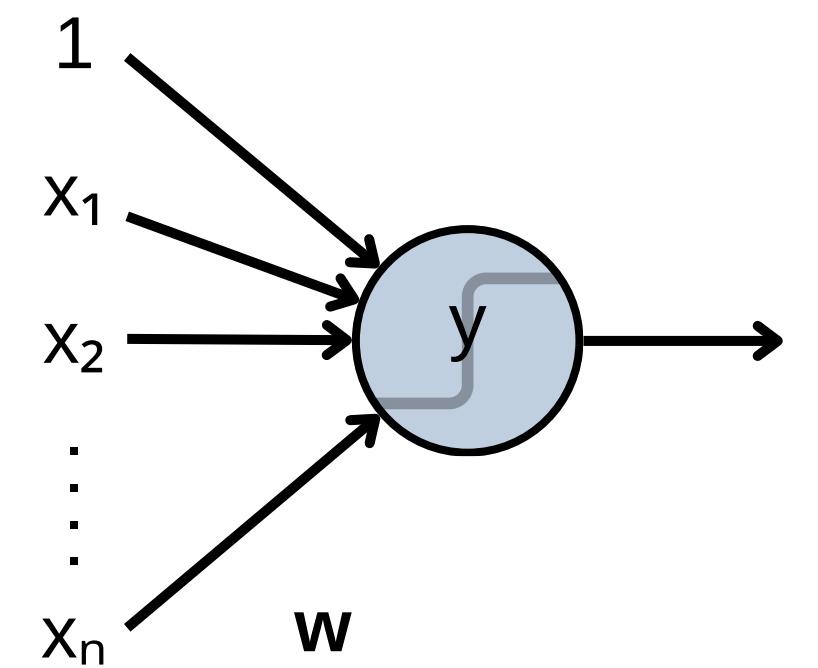
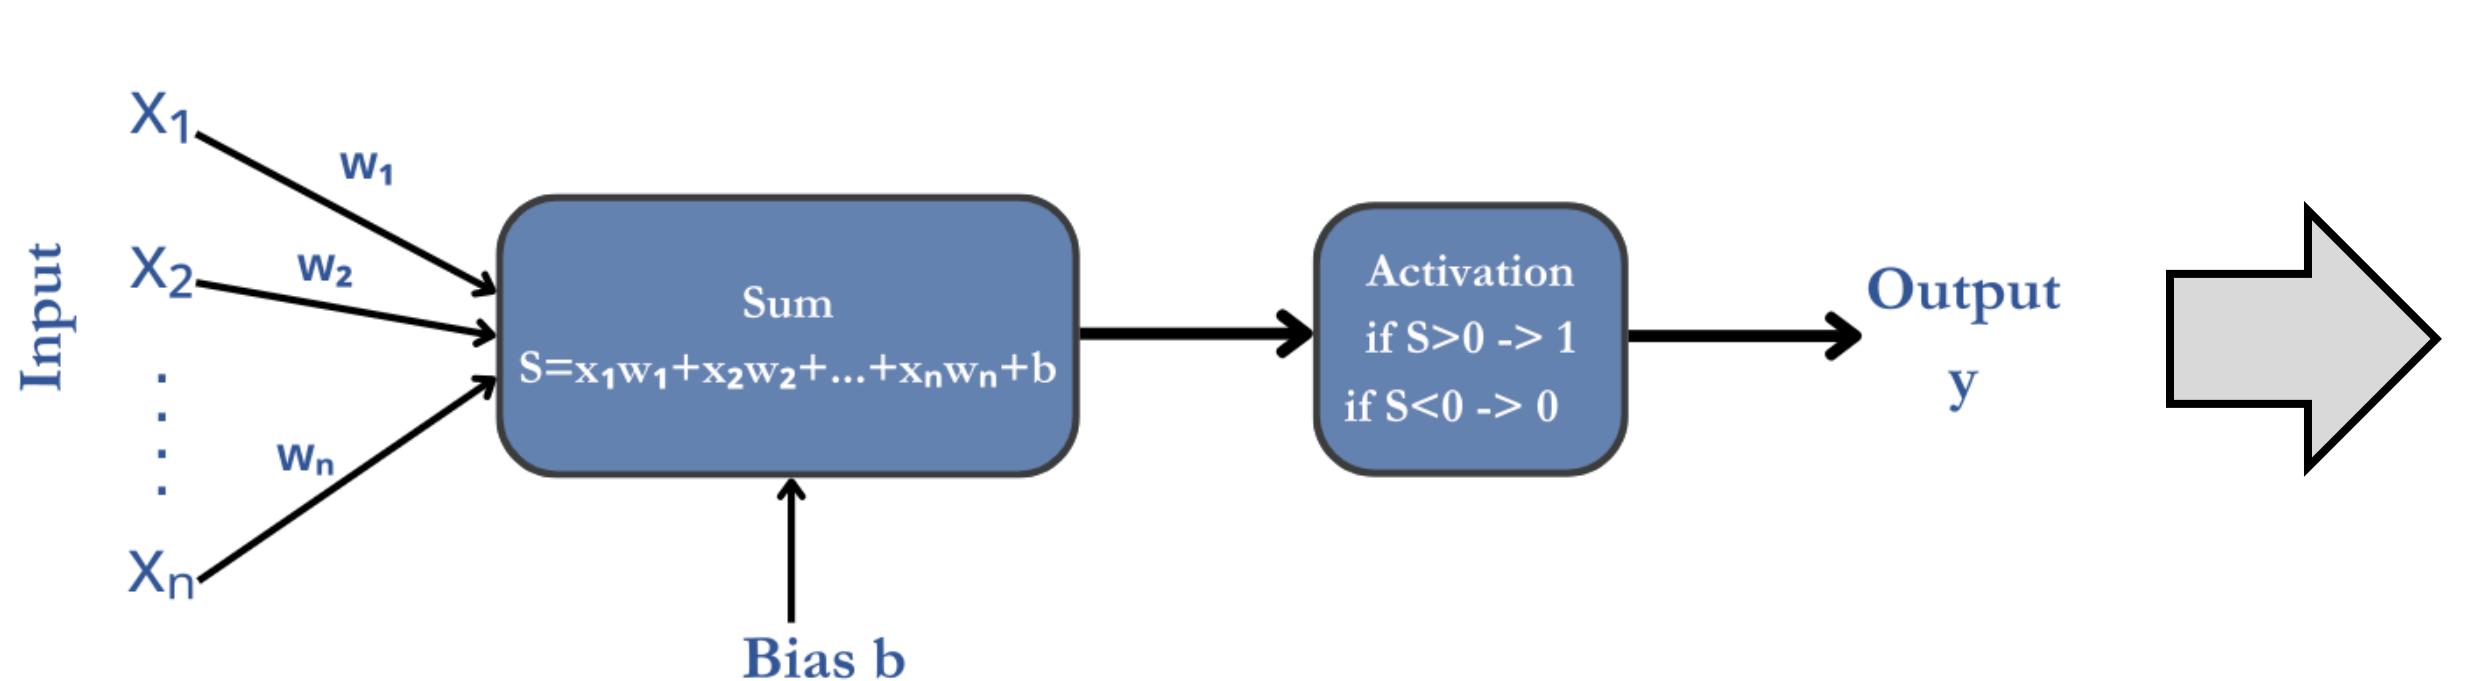
# Shallow Neural Networks



# Another Representation

In the following we will use a more simple representation for the perceptron

- we combine the weights and the bias in the same vector
  - $w = (b, w_1, w_2, \dots, w_n)$
- we add a dummy input that is always 1 and that gets multiplied by the bias
- we write on the neuron the name of its output and we plot its activation on it

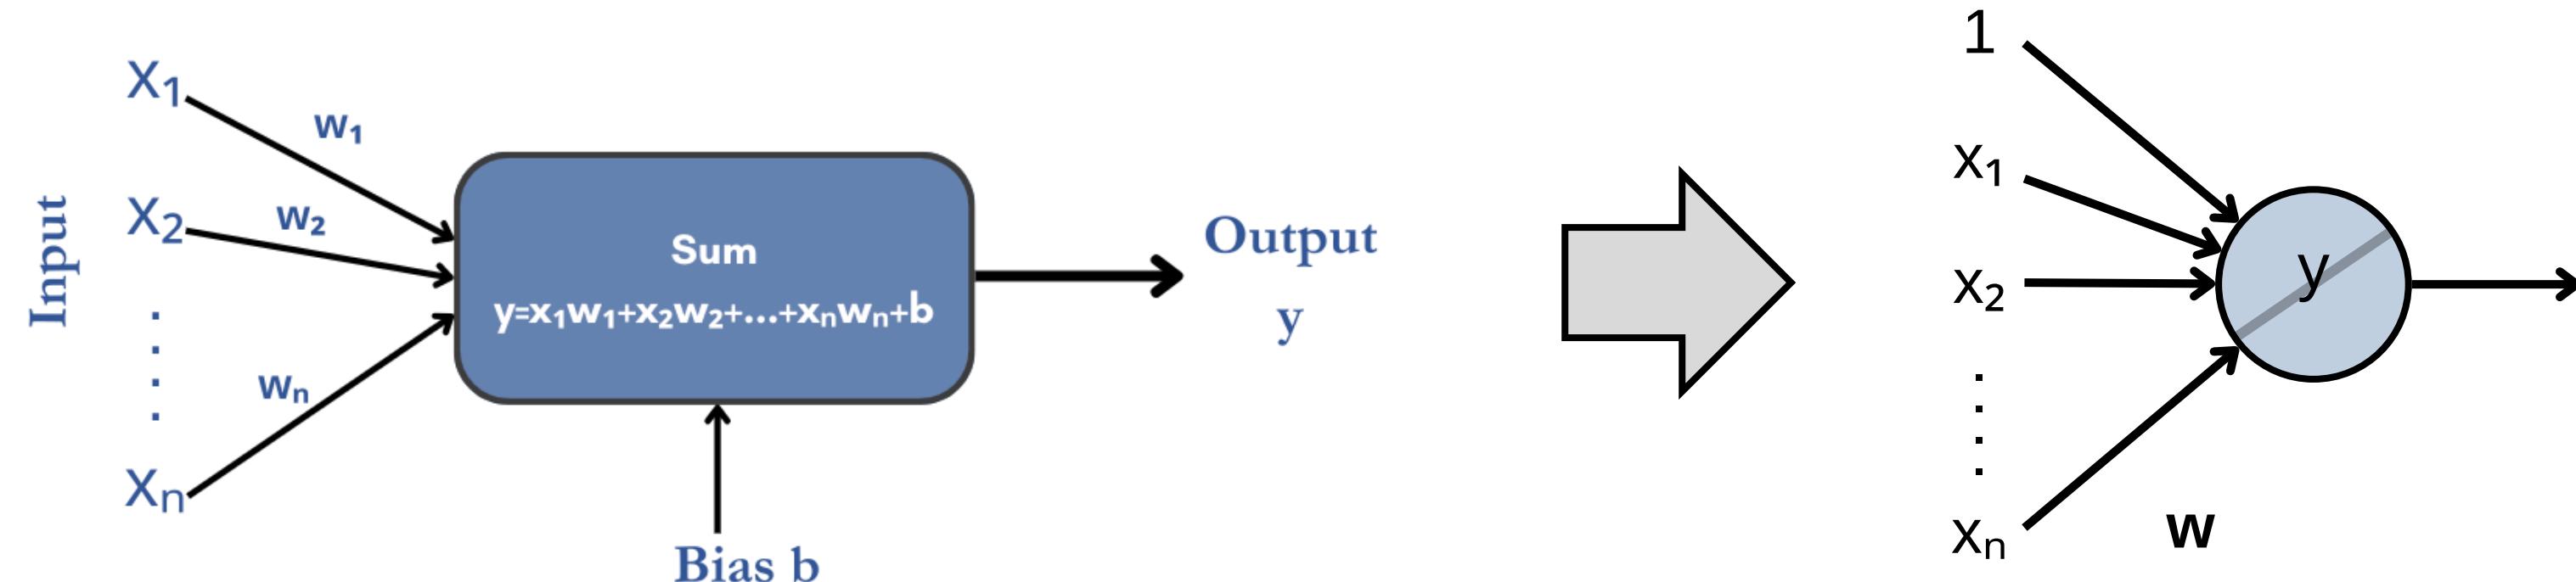




# Another Representation

In the following we will use a more simple representation for the perceptron

- we combine the weights and the bias in the same vector
  - $w = (b, w_1, w_2, \dots, w_n)$
- we add a dummy input that is always 1 and that gets multiplied by the bias
- we write on the neuron the name of its output and we plot its activation on it





# Activation Functions

## Step Function

- Used as output for the classification task
  - $a(x) = 1$  if  $x > 0$
  - $a(x) = 0$  if  $x < 0$

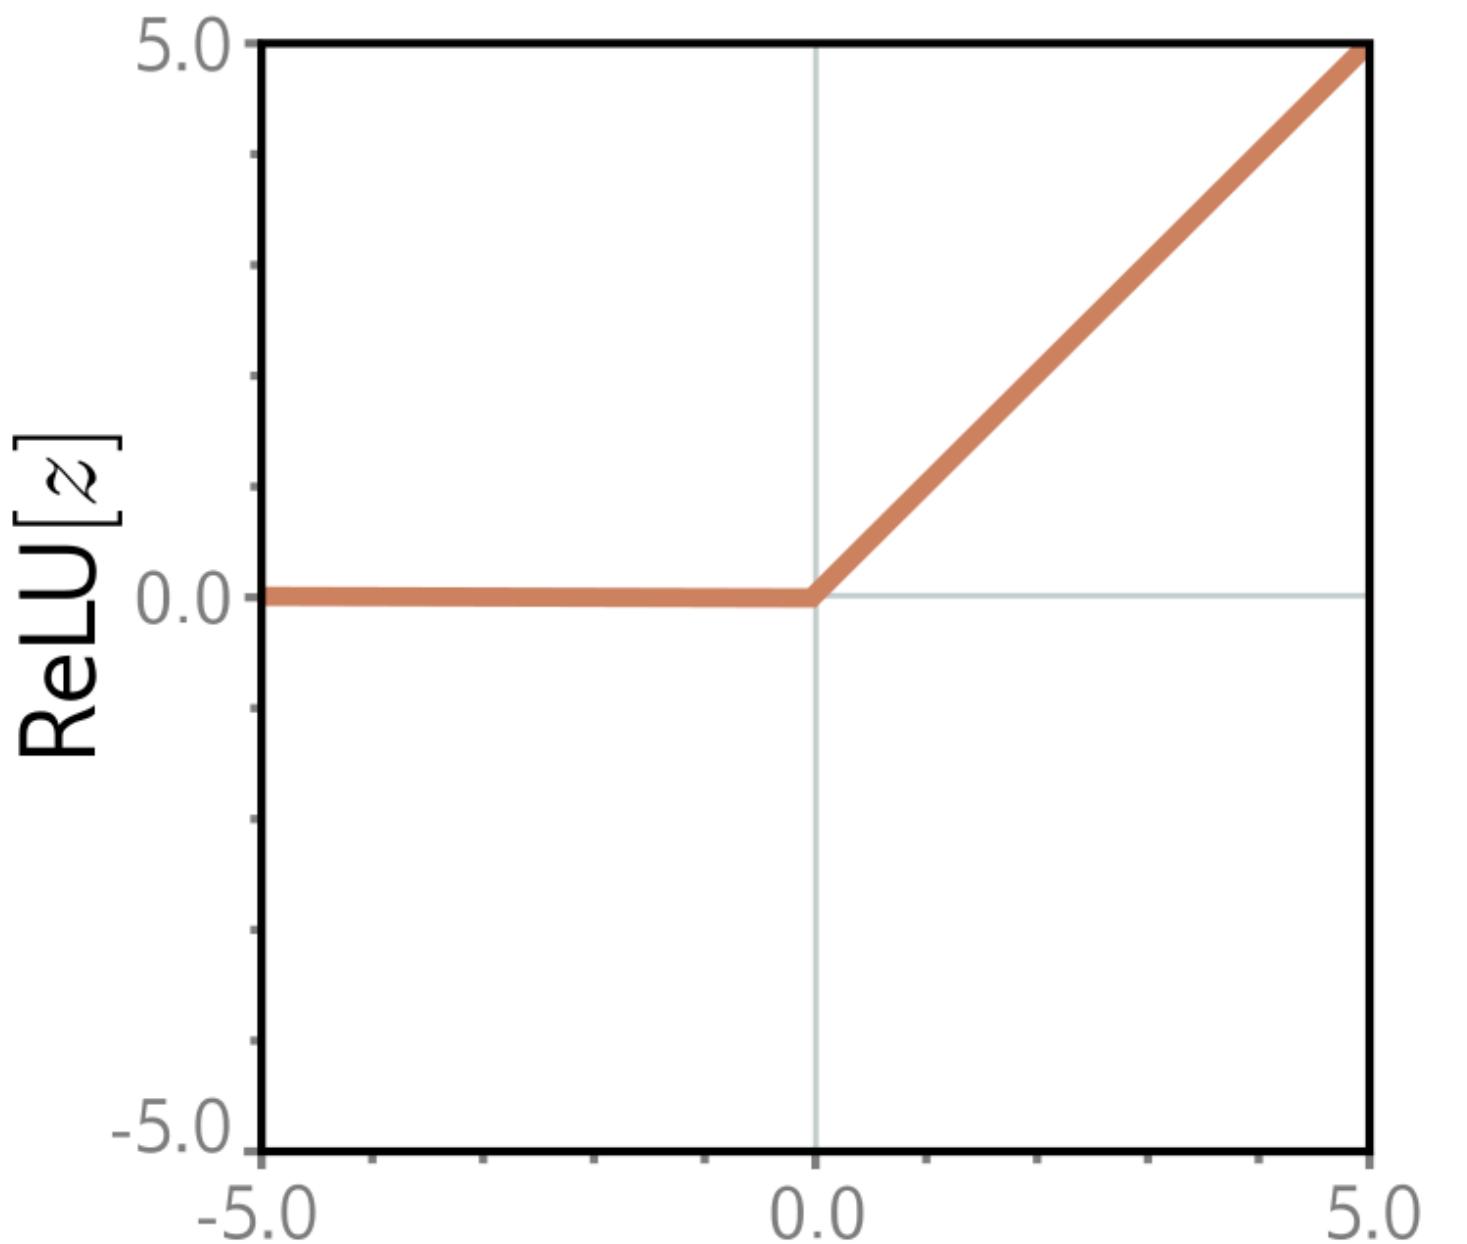
## Linear Activation

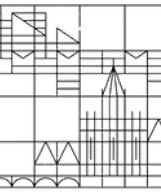
- Used as output for the regression task
  - $a(x) = x$

The Rectified Linear Unit (ReLU) is another example

## ReLU

- Used in hidden layers of deep neural networks
  - $a(x) = x$  if  $x > 0$
  - $a(x) = 0$  if  $x < 0$



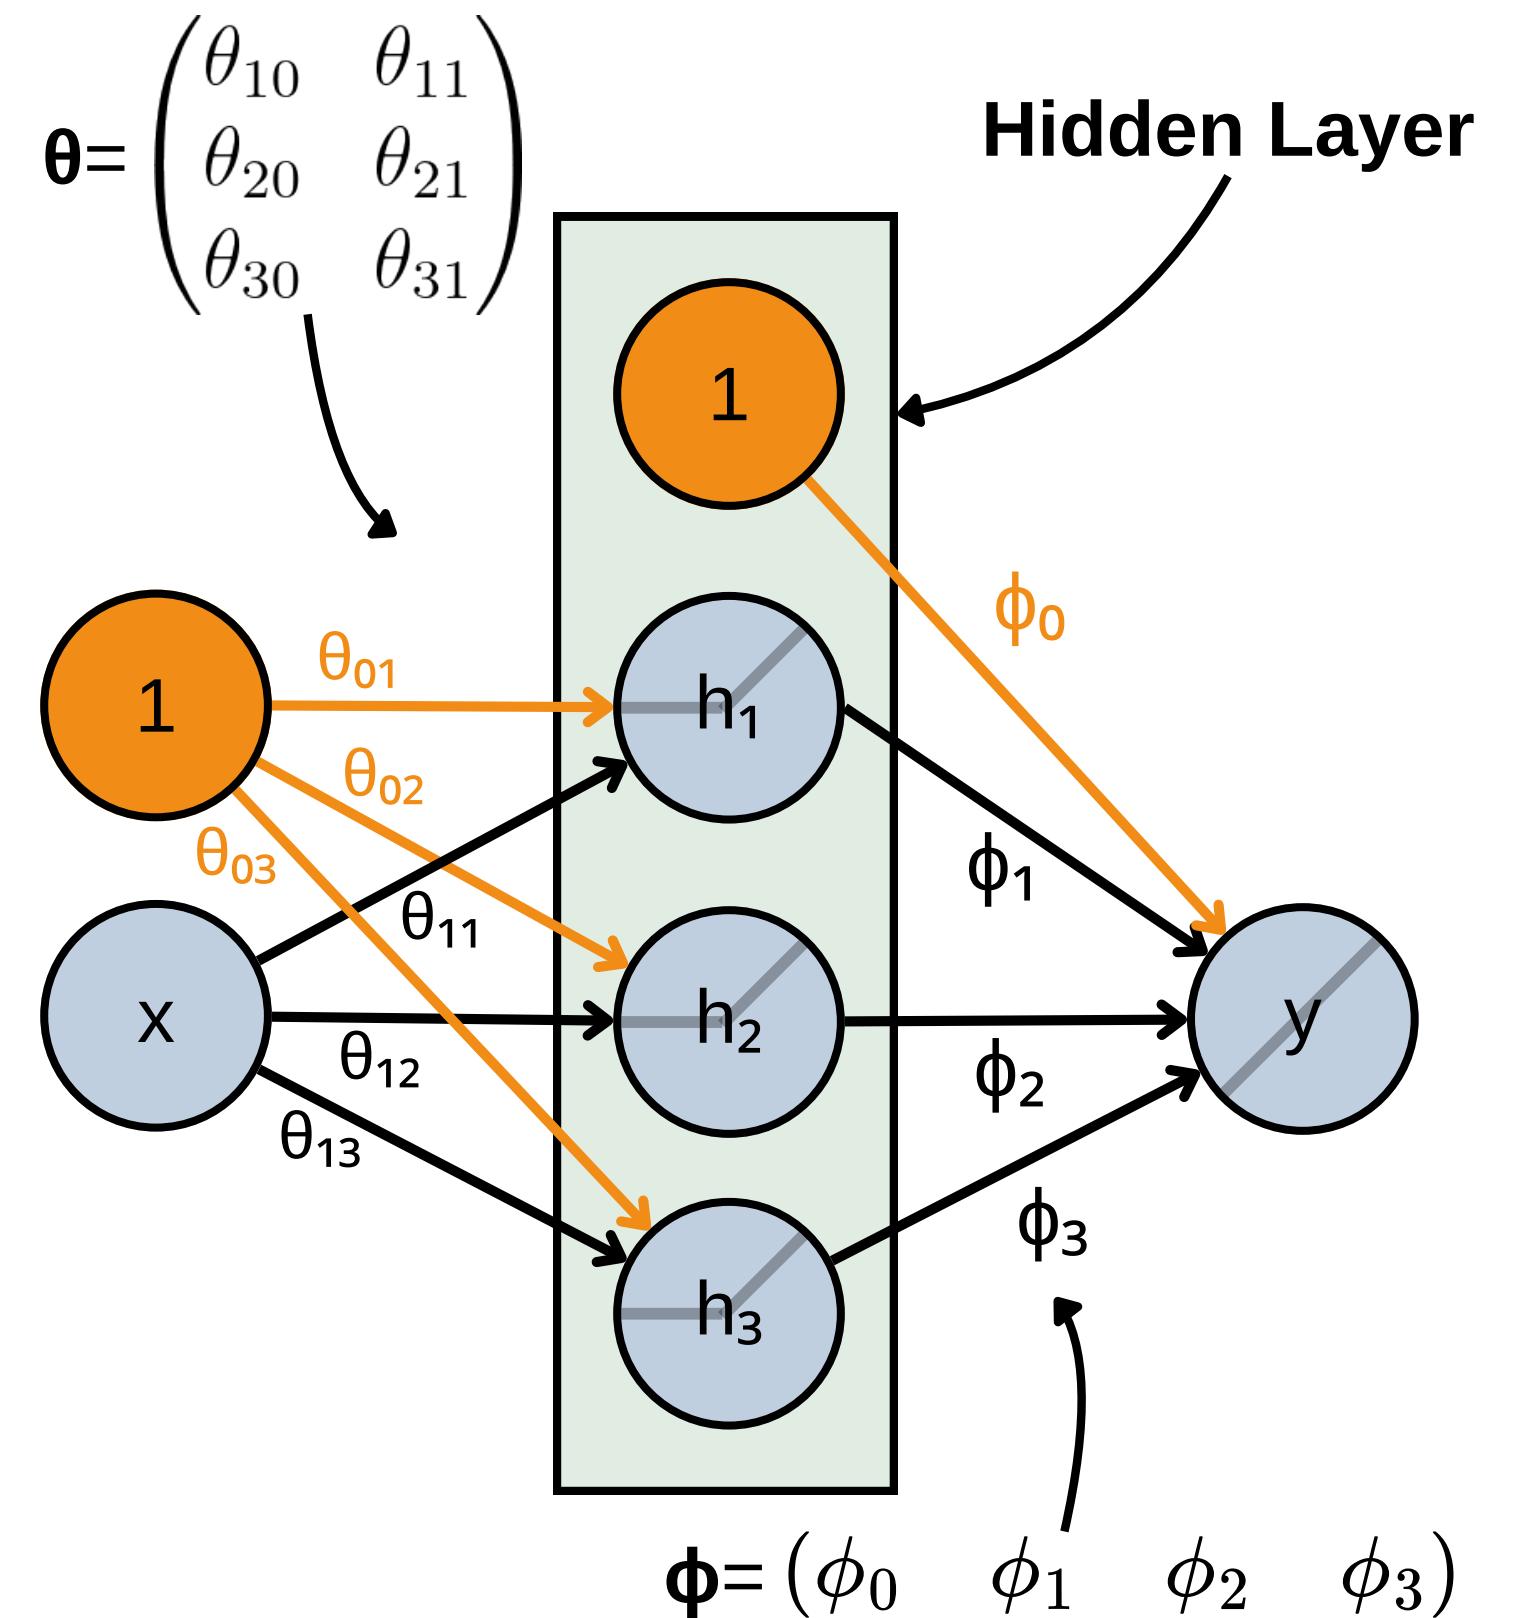


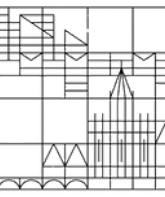
# Combining Perceptrons

We consider again the simple regression problem with a single input  $x$  and output  $y$

- we can apply more than one single perceptron to the input (and dummy)
- each of these perceptrons will produce a different output  $h_i$
- we can then use these outputs as input for another perceptron that produce the output  $y$

In this way we are adding an **hidden layer** to the neural network





# Shallow Neural Networks

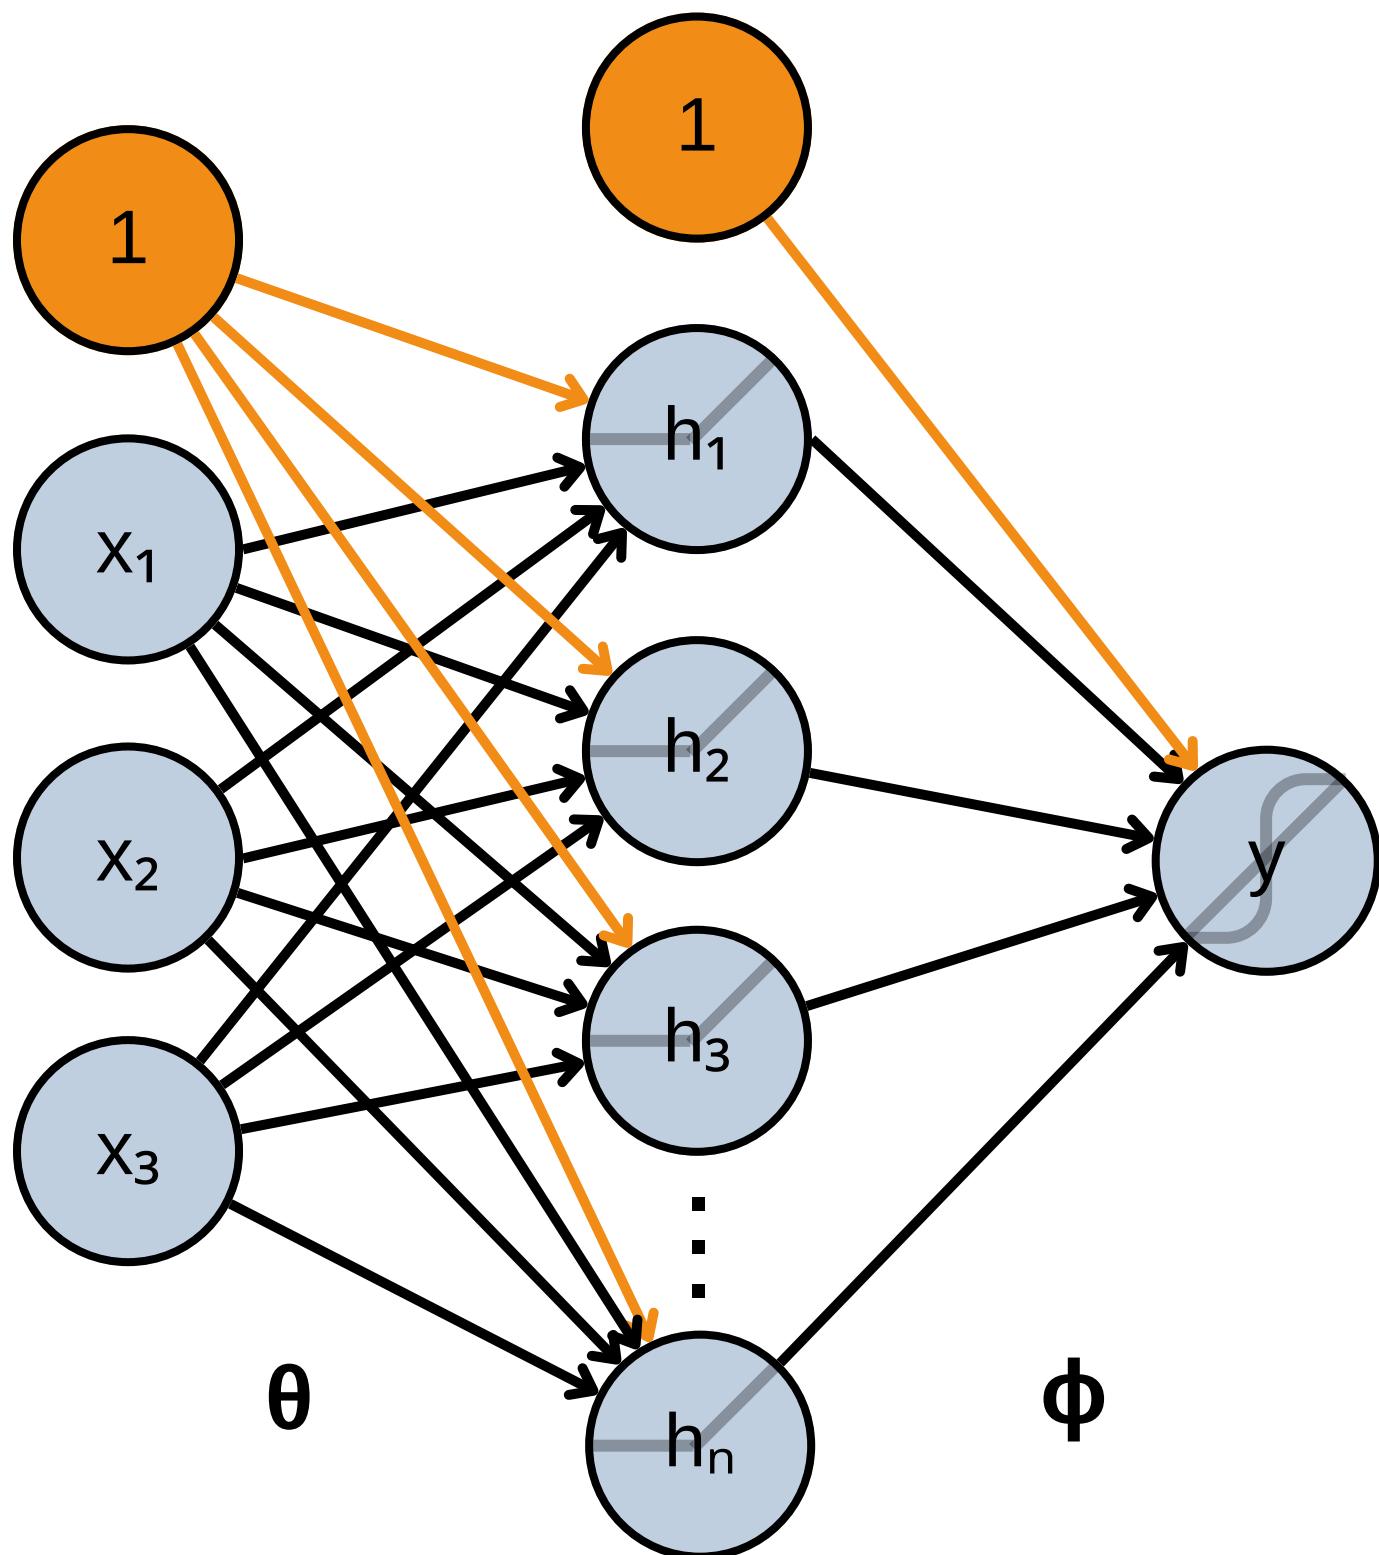
More generally we can have neural networks with

- as many inputs as we want
- as many hidden neurons as we want

The parameters of this neural network will be contained in two weight matrices

- $\theta$  connecting the input to the hidden layer
- $\phi$  connecting the hidden layer to the output

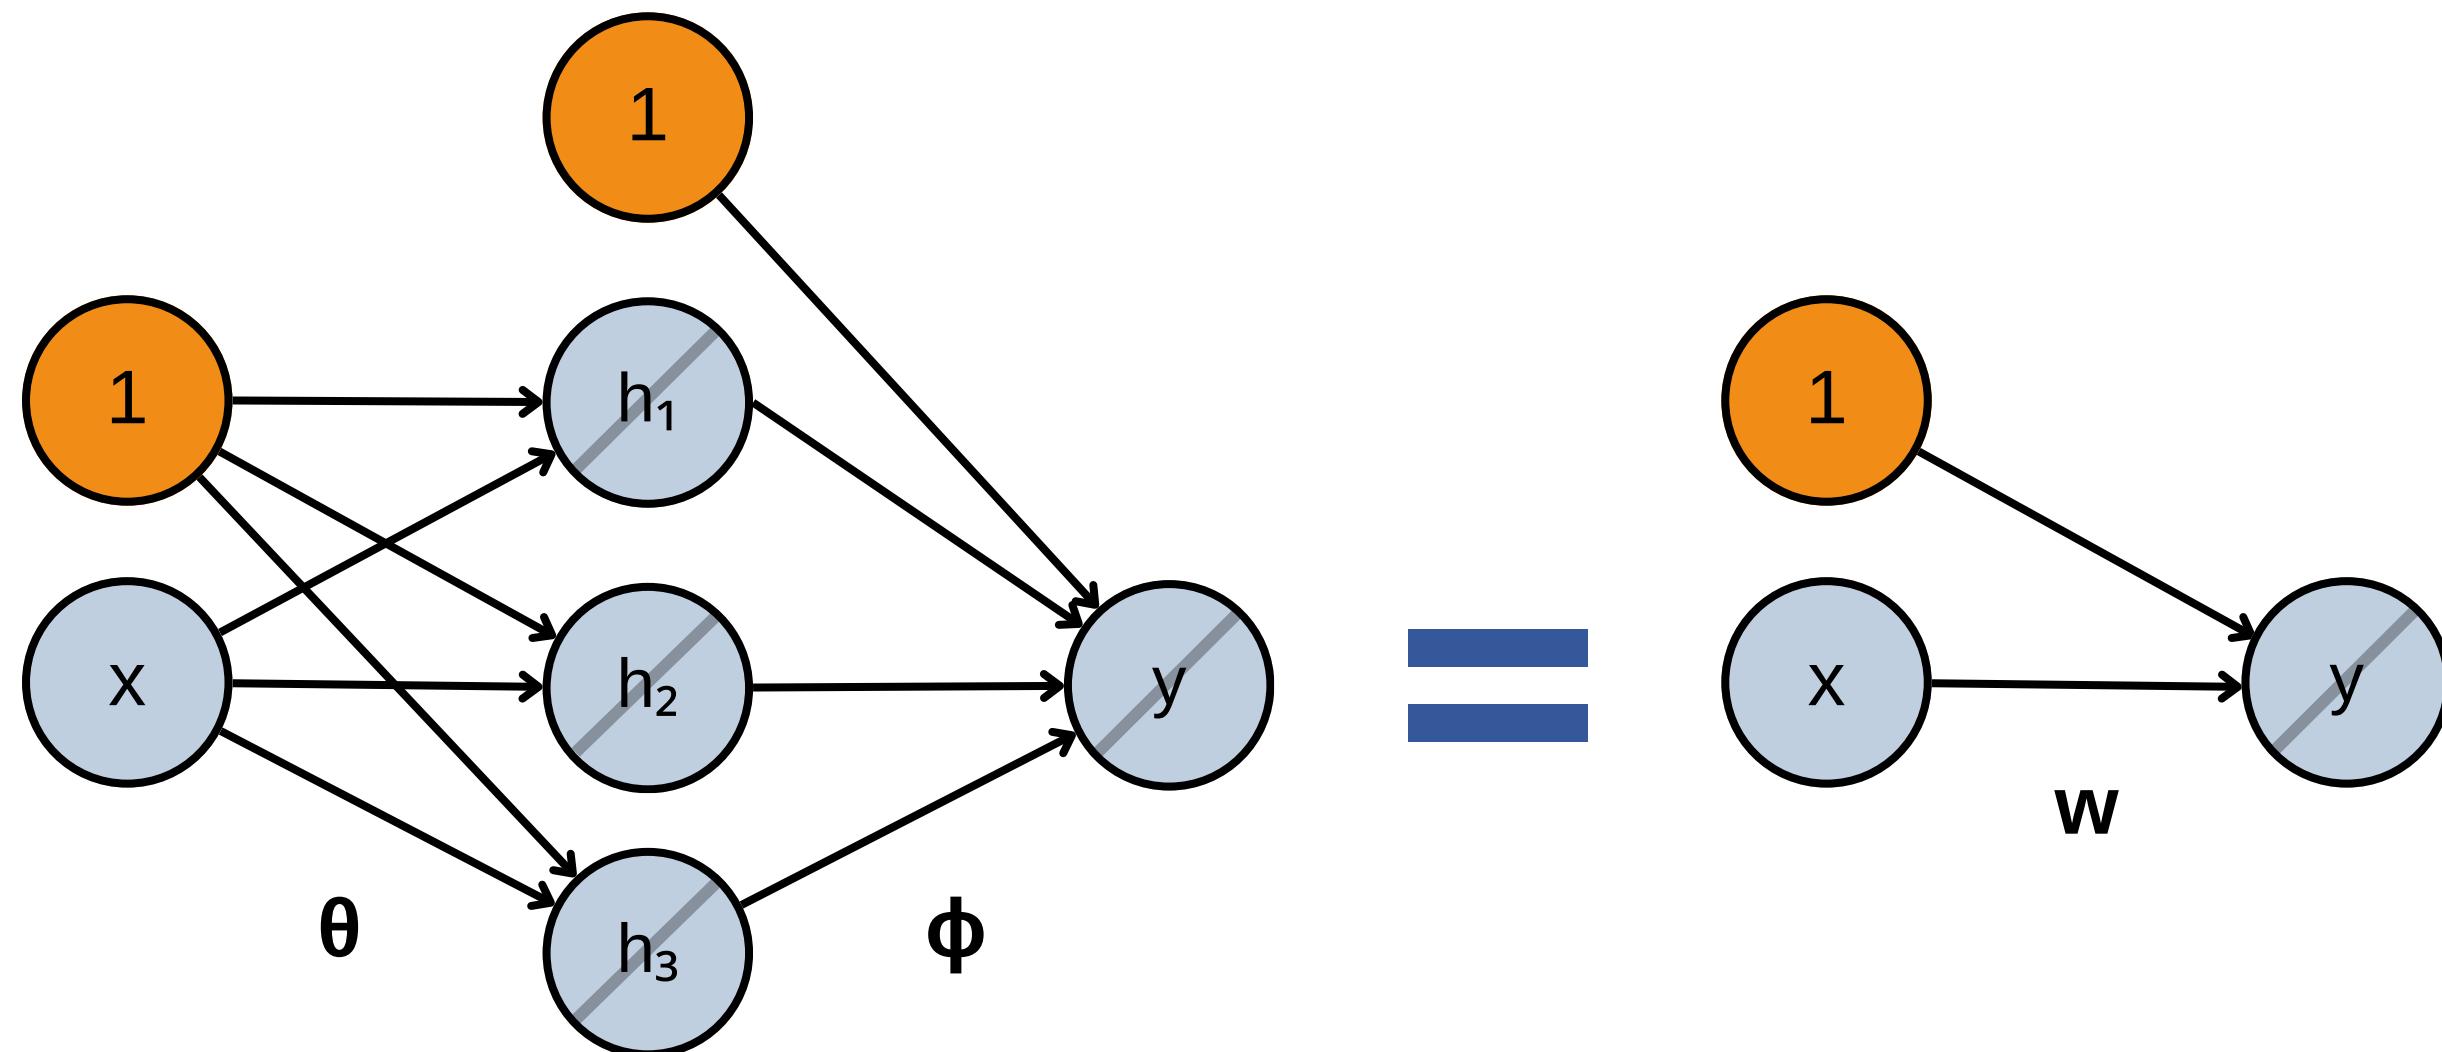
This type of neural network with a single hidden layer is called **Shallow Neural Network**

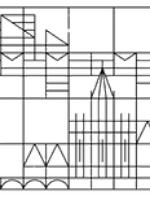




# Why is the ReLU Important?

Non-linear activation functions like the ReLU are crucial in Deep Learning. A shallow neural network with linear activation functions is equivalent to a simple perceptron





# What's Next?

The perceptron can only solve liner problem

- most real life problems are much more complex
- this lead to the AI Winter

We have just introduced shallow neural networks

- if we use non-linear activation functions like the ReLU this is different with respect to a simple perceptron

We still have to understand some things

- can shallow neural networks solve non-linear problems?
- how can we train a shallow neural network?



# What's Next?

Tomorrow room G420

- introduction to google colab and GPU server
- basic machine learning concepts
- fill the form to get access to the GitHub of the course

[https://docs.google.com/forms/d/e/1FAIpQLSc9bKHplUFxv\\_jxfY20OYmA0OrjilCcAaCMEICREQtA0t9Q2w/viewform?usp=sharing](https://docs.google.com/forms/d/e/1FAIpQLSc9bKHplUFxv_jxfY20OYmA0OrjilCcAaCMEICREQtA0t9Q2w/viewform?usp=sharing)

Next week

- on Wednesday we answer to the open questions and we introduce the first Deep Neural network, the Multilayer Perceptron
- on Thursday we will implement our first neural network using PyTorch



# Summary

## Basic Concepts and Notation

We introduced the main machine learning concepts like supervised vs unsupervised learning, classification and regression, the loss function

## The Perceptron

The perceptron is the first artificial neural network. It consists in a weighted sum and an activation and allows to perform automatic classification/regression

## Limits of the Perceptron

The perceptron can only solve linear problems: in the case of classification it draws a linear decision boundary, while in the case of regression it can only perform a linear regression

## Shallow Neural Networks

The output of a perceptron can be fed into another perceptron, leading to a shallow neural network. The hidden layer must have non-linear activation functions like the ReLU